

# Photography Made Easy

Sylvain Paris, Adobe



Photo before retouching



After retouching

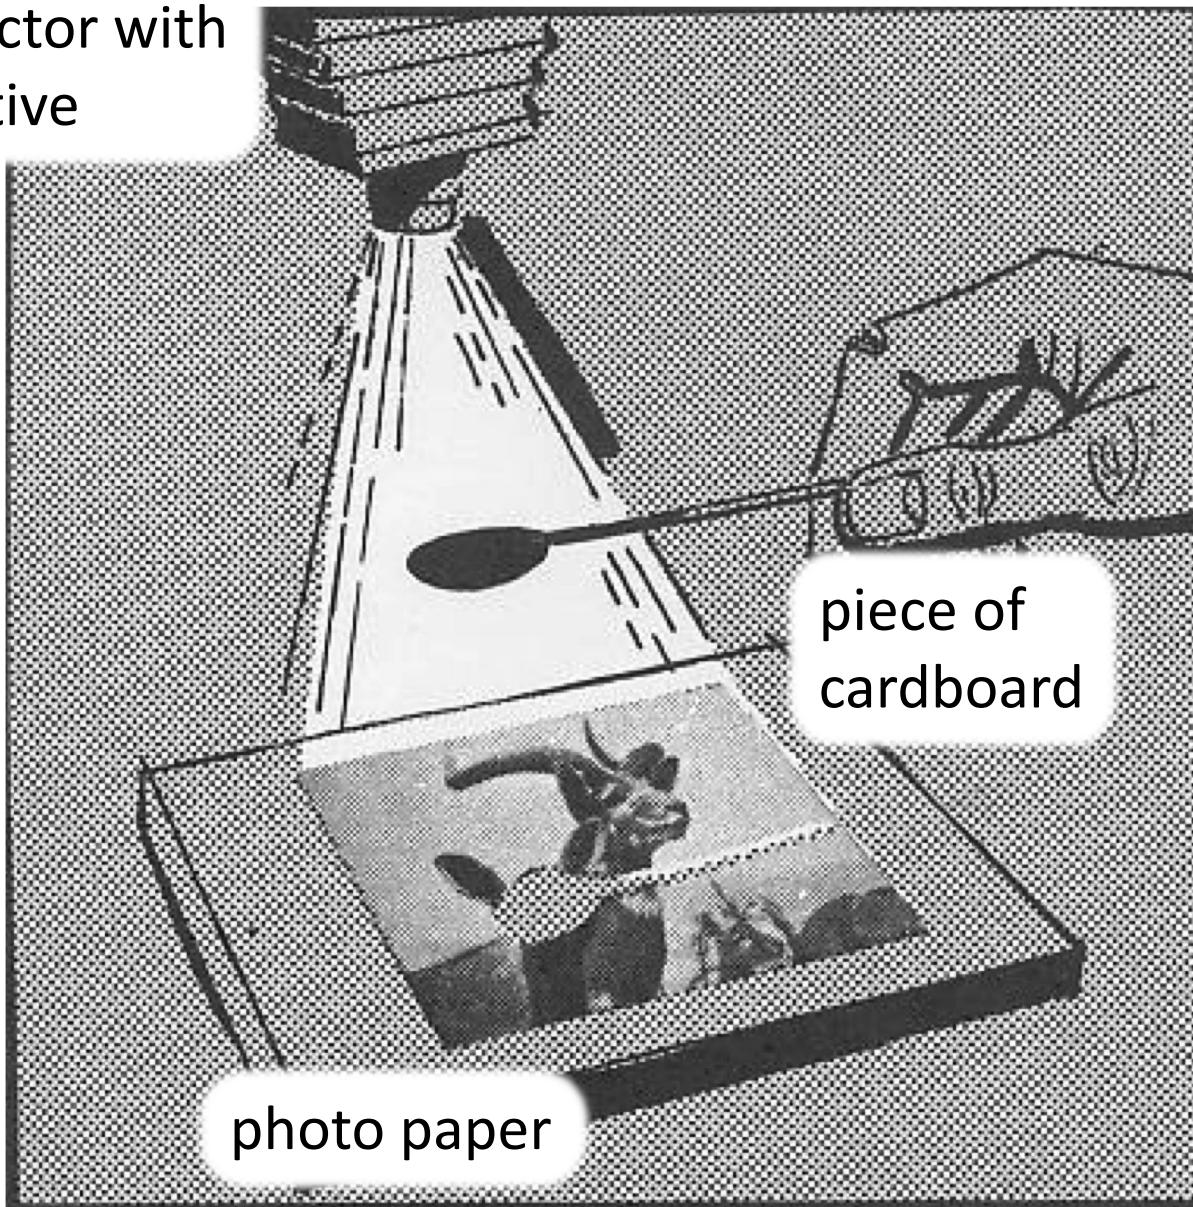


Photo before retouching

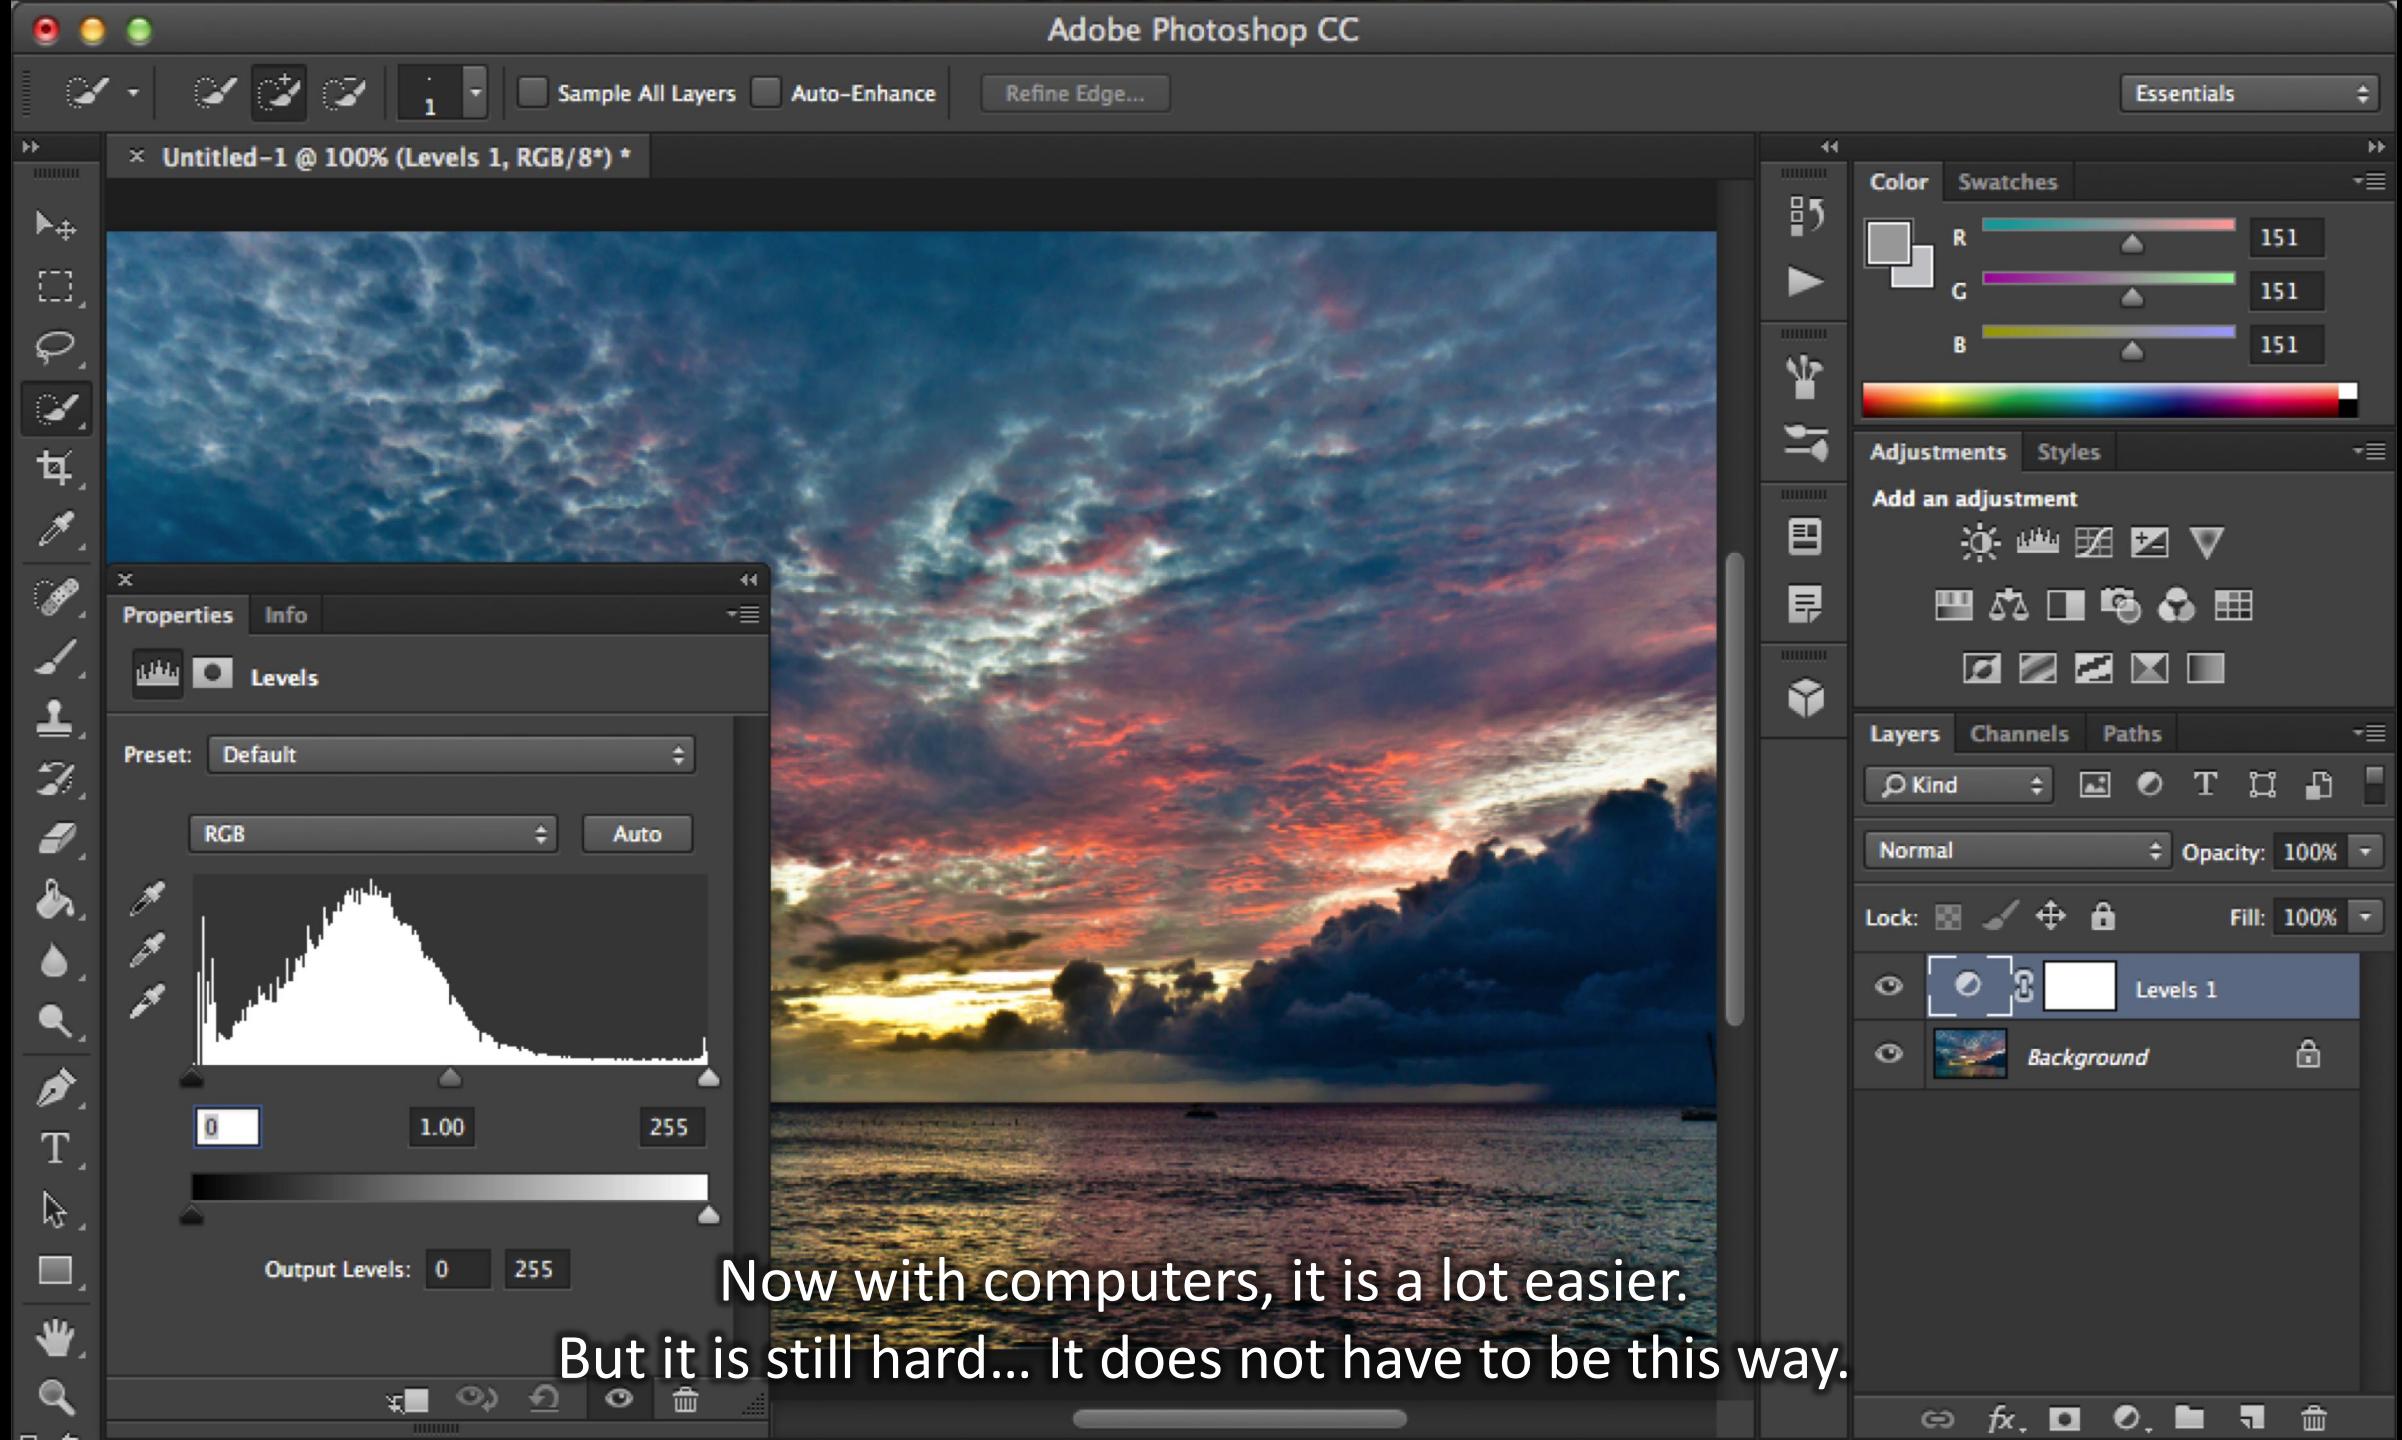


After retouching

projector with  
negative



Printing photos used to be hard



Now with computers, it is a lot easier.  
But it is still hard... It does not have to be this way.

Many photos can become great after retouching.

I want to make it easy.

# How to Make Photo Retouching Easier

- Some people know how to do it: photographers and artists
- There are plenty of examples of good photos available on the web
- Our strategy
  1. Develop algorithms that transfer the statistics of good photos.
  2. Learn how to retouch from examples of good photos

# Photographic Style Transfer

*SIGGRAPH 2006*

Make this photo



look like



# Tonal Aspects of Look: Global Contrast



Ansel Adams

**High global contrast**



Kenro Izu

**Low global contrast**

# Tonal Aspects of Look: Local Contrast



Ansel Adams

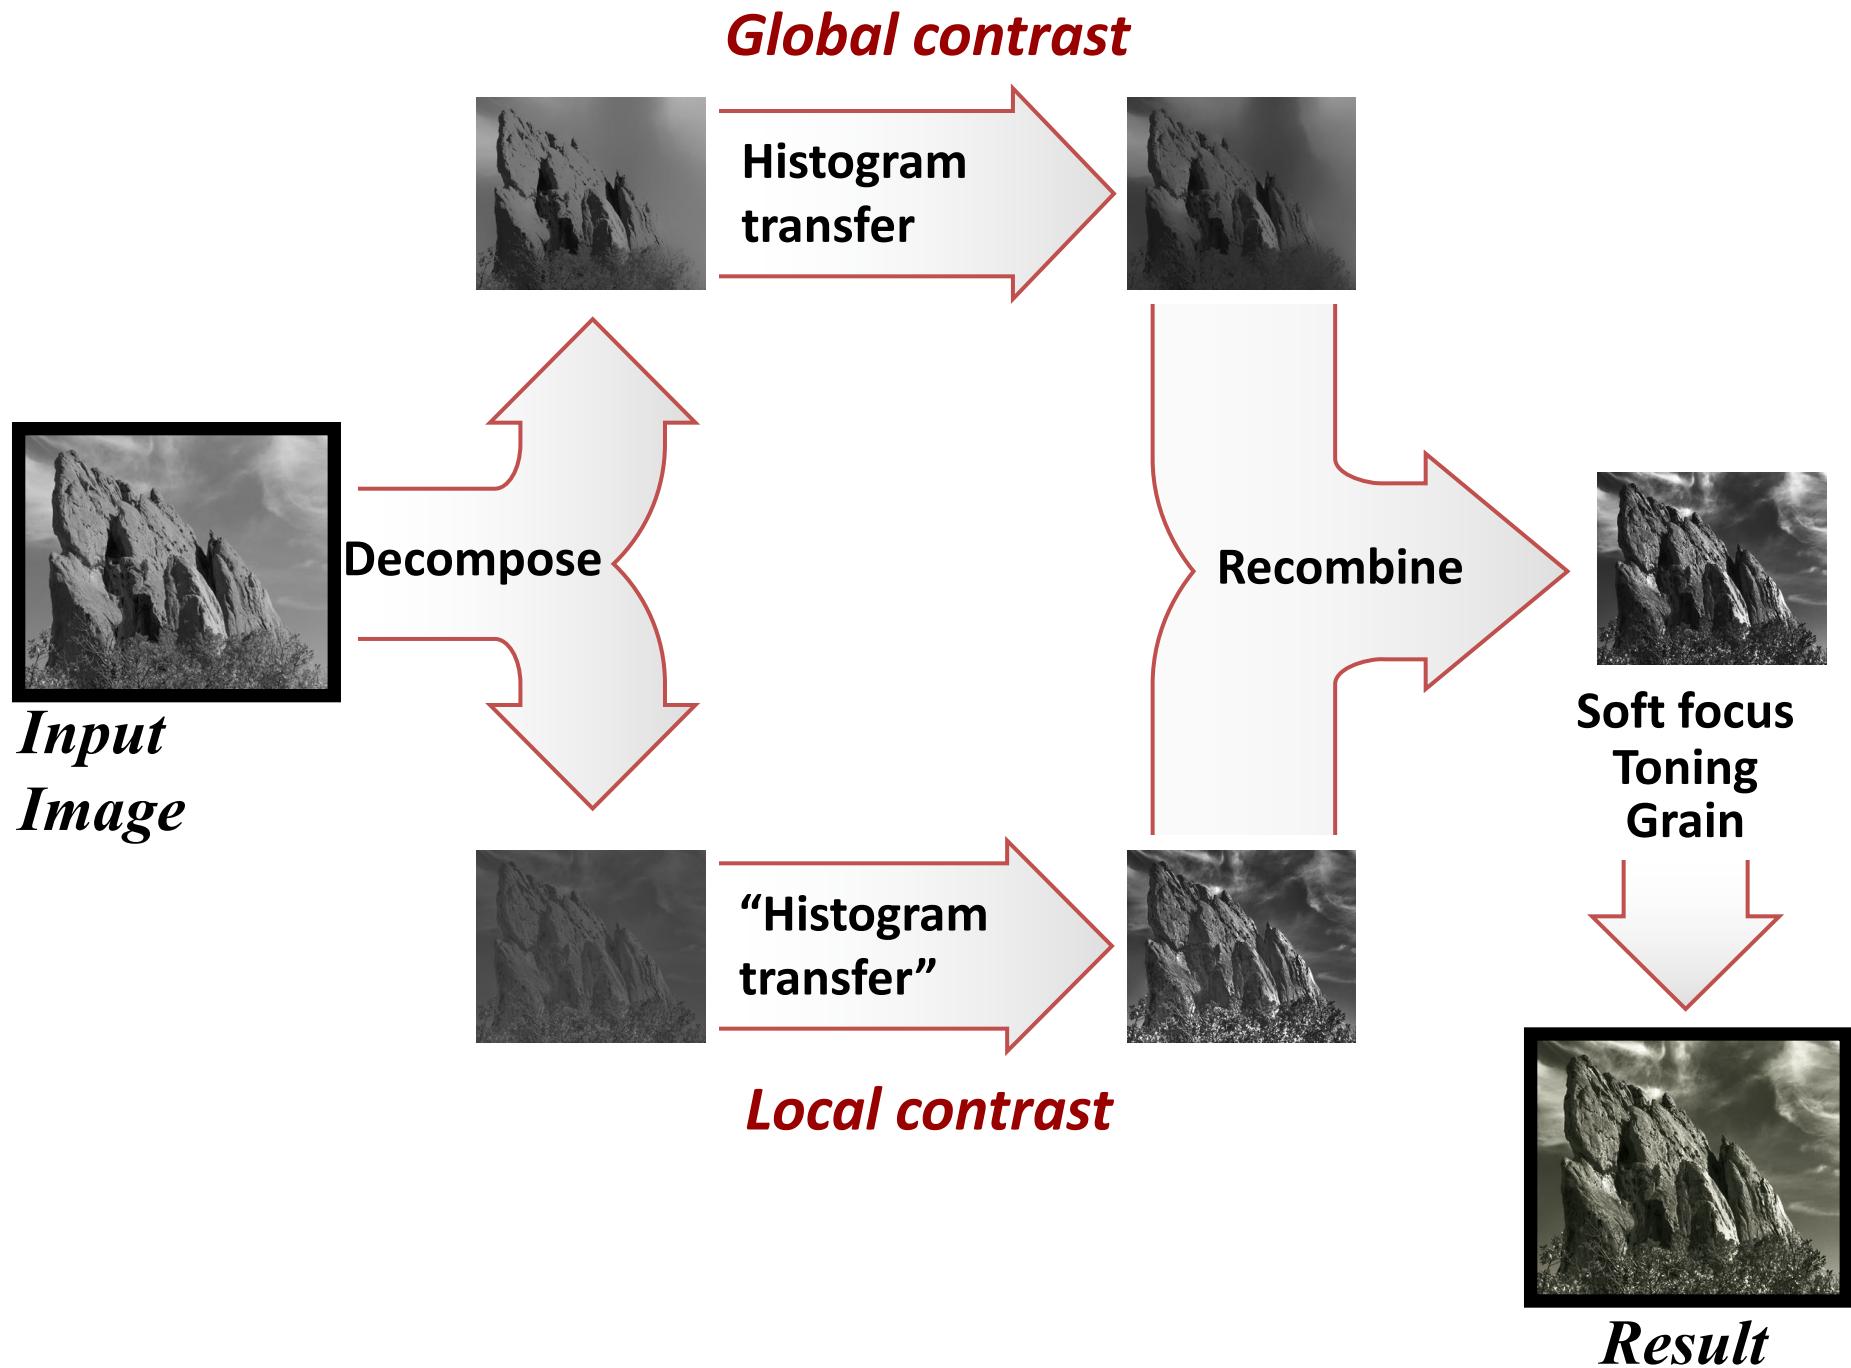


Kenro Izu

**Variable amount of texture**

**Texture everywhere**

# Pipeline

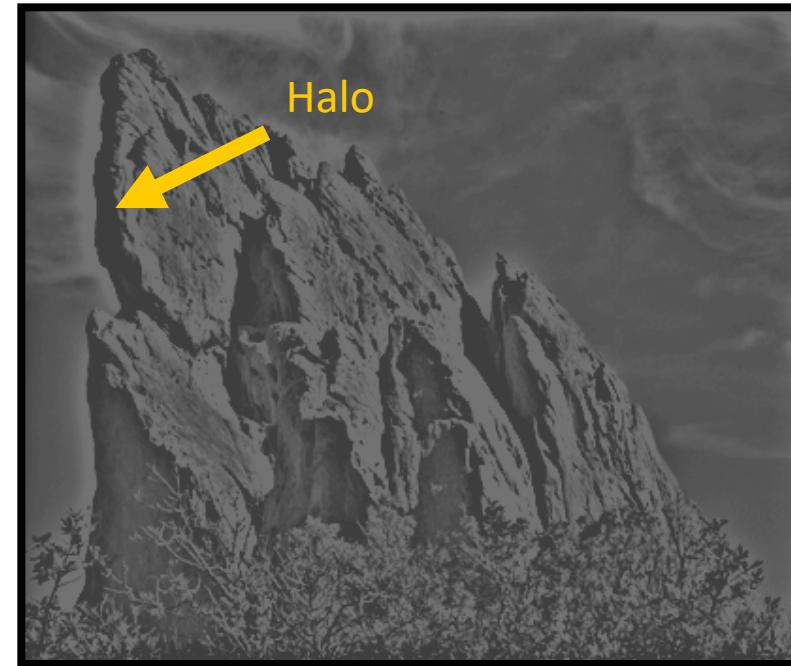


# Naïve Decomposition: Low vs. High Frequency

- Problem: introduce blur & halos



Low frequency  
*Global contrast*



High frequency  
*Local contrast*

# Edge-Preserving Decomposition: Bilateral Filter



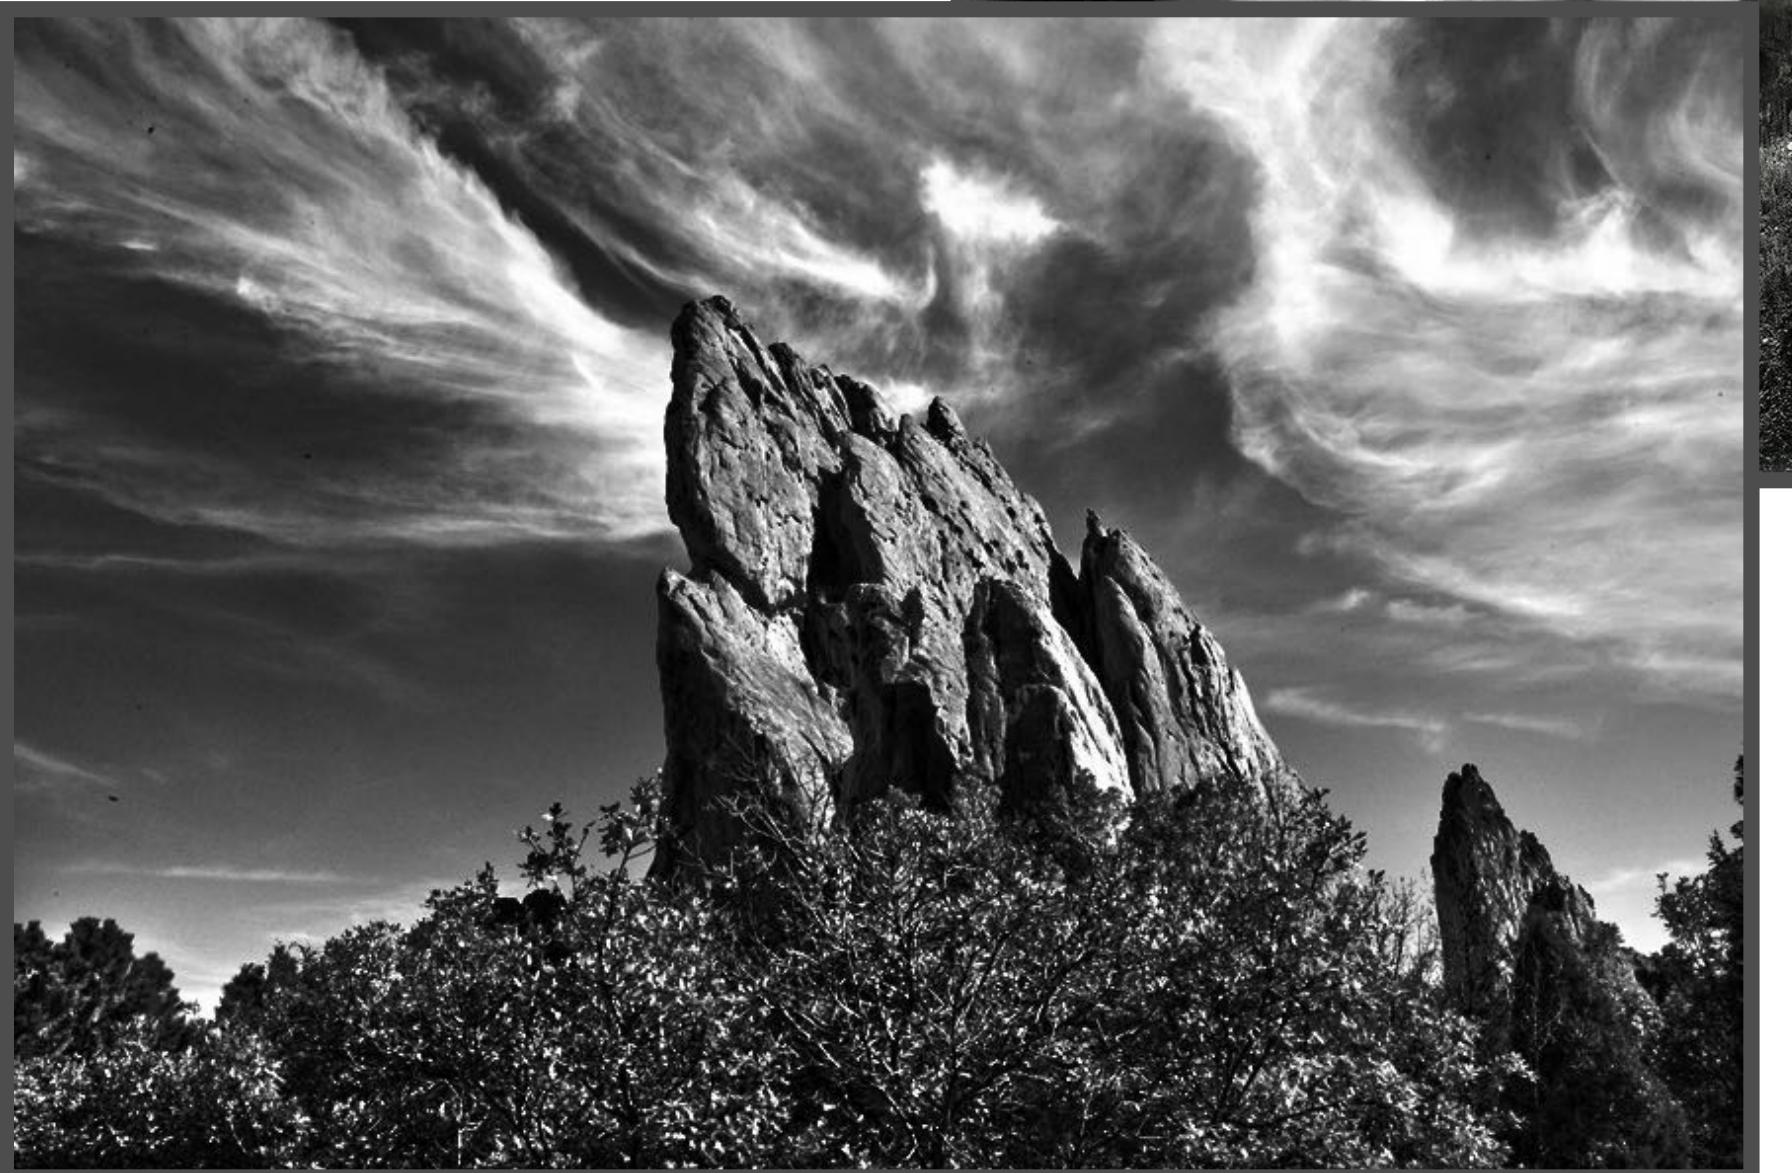
Bilateral filter output  
*Global contrast*



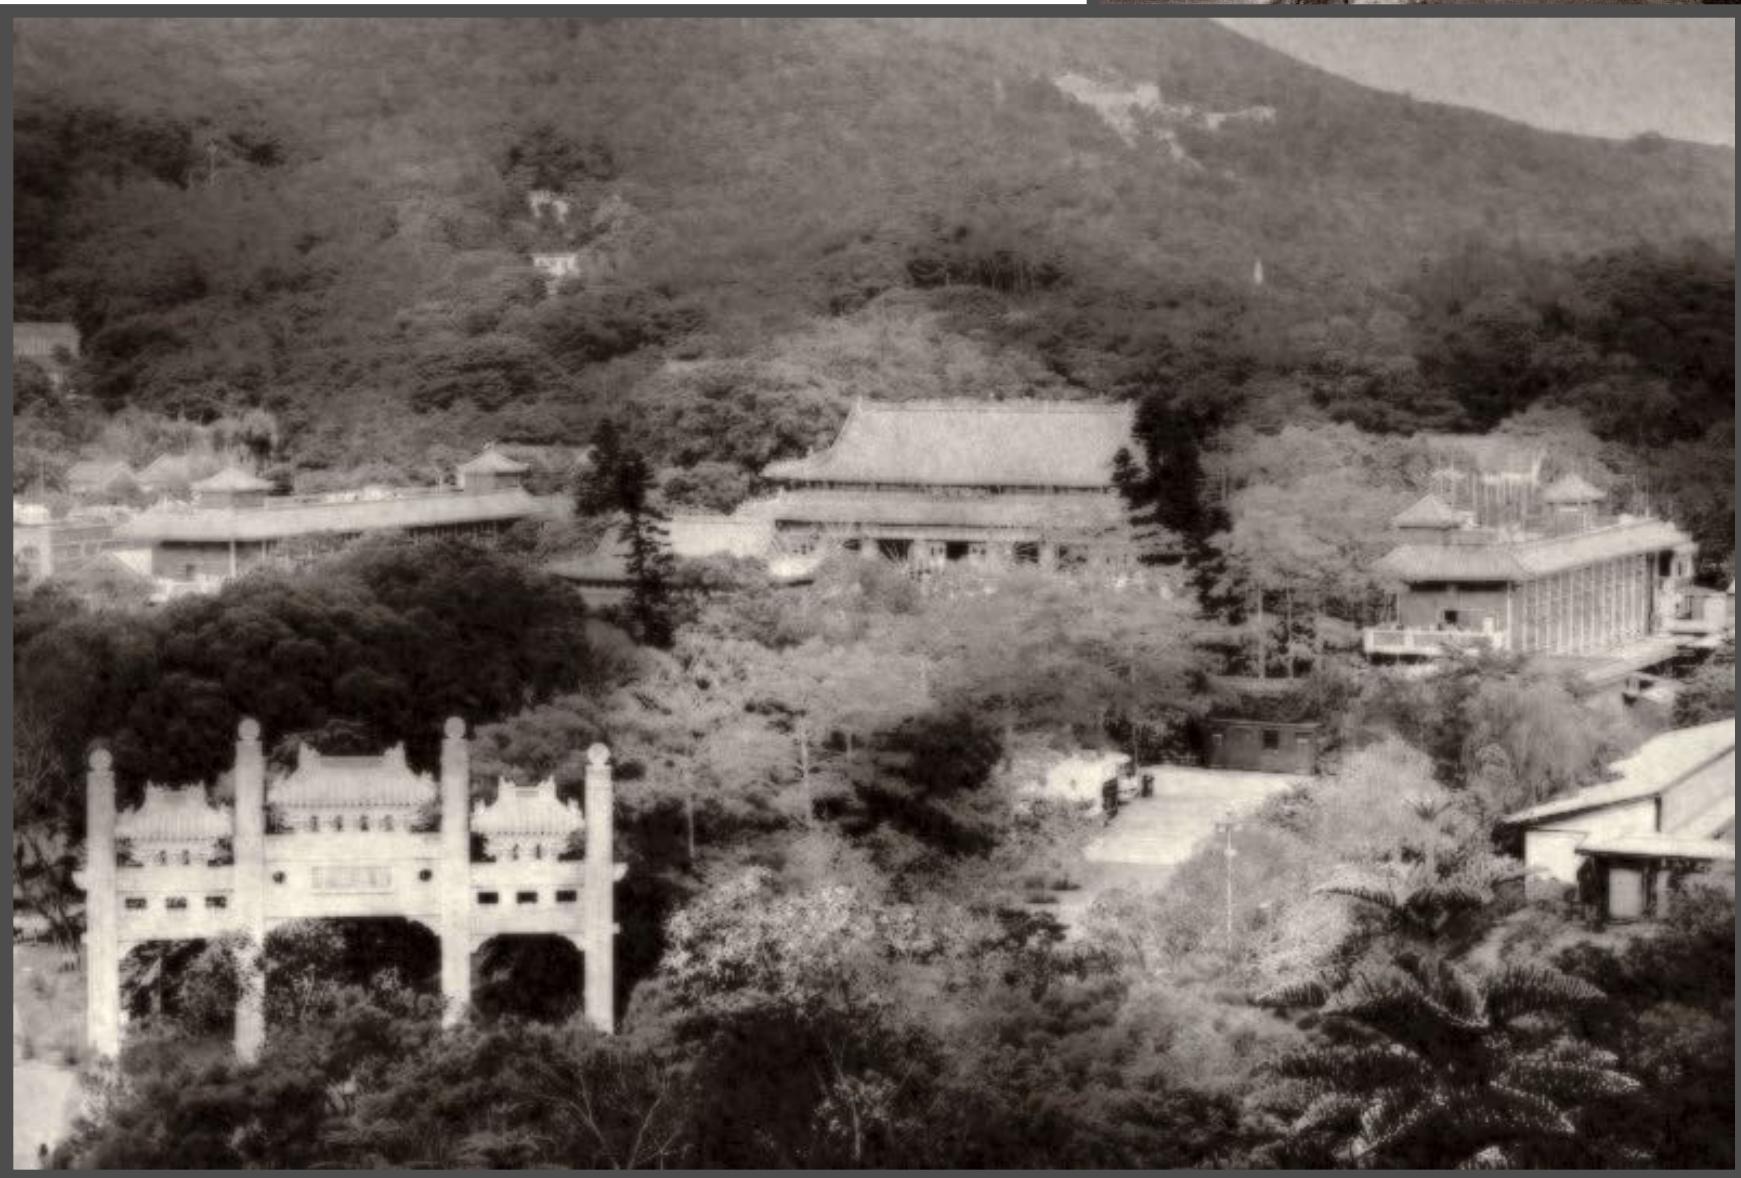
Bilateral filter residual  
*Local contrast*

**Result**

**Model**



# Result



# Photoshop Demo

# Local Laplacian Filters

*A Better Decomposition*

SIGGRAPH 2011

# Background on Gaussian Pyramids

- Resolution halved at each level using Gaussian kernel



level 0



level 1



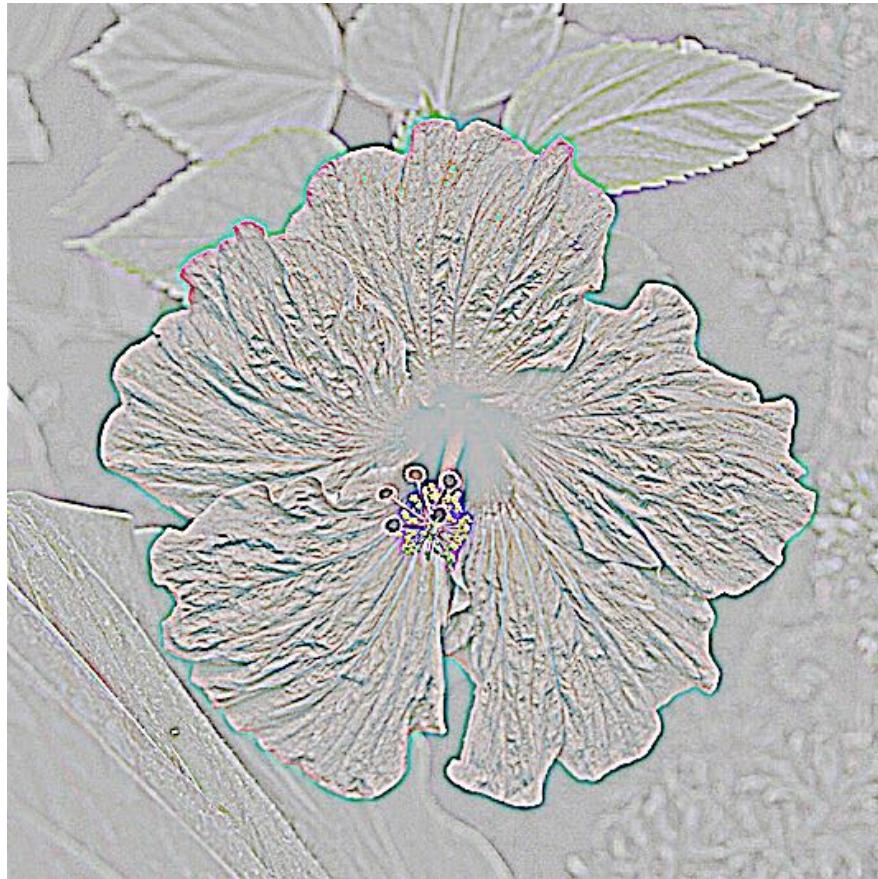
level 2



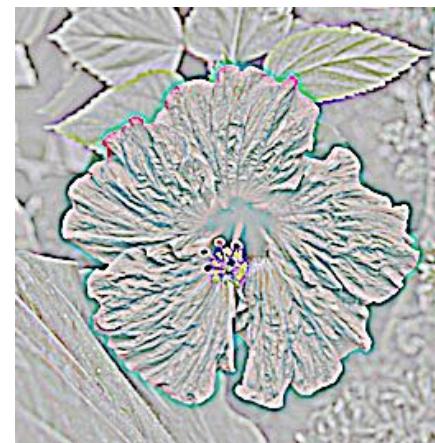
level 3  
(residual)

# Background on Laplacian Pyramids

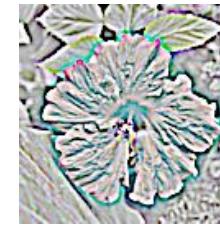
- Difference between adjacent Gaussian levels



level 0



level 1



level 2



level 3  
(residual)

# Pros and Cons of Pyramids

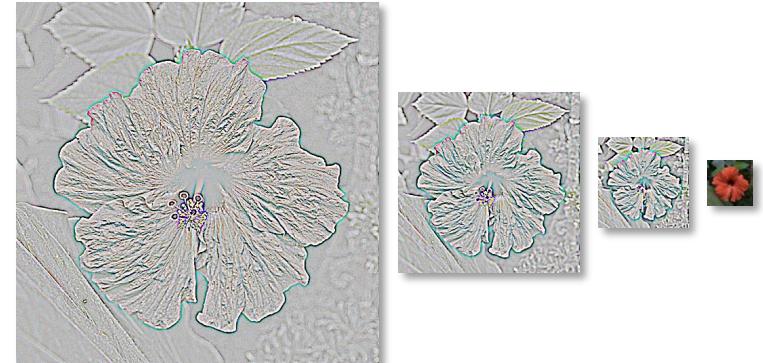
- ☺ Useful for compression [Burt 83],  
texture synthesis [Heeger 95],  
harmonization [Sunkavalli 10]...
  
- ☹ Believed to be unsuitable for  
edge-aware processing
  - Use isotropic & spatially invariant kernels
  - But edges are anisotropic & well located
  - “*Manipulating pyramids generate halos*”
- We show otherwise.



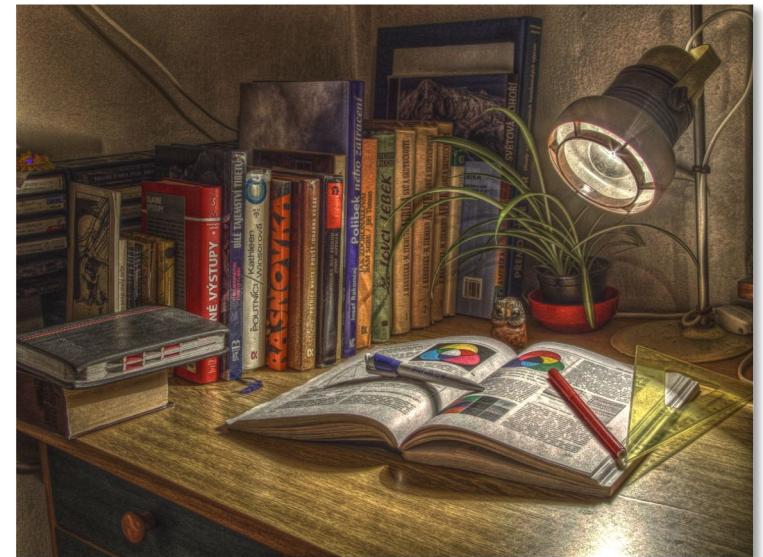
*boosting high-frequency bands*

# Our Contributions

- **Edge-aware editing with Laplacian pyramids**
  - We use a classical multi-scale representation



- **Robustness:** strong effects, no artifacts
  - We achieve extreme enhancements where other methods fail



# Our Strategy: Local Adaptation

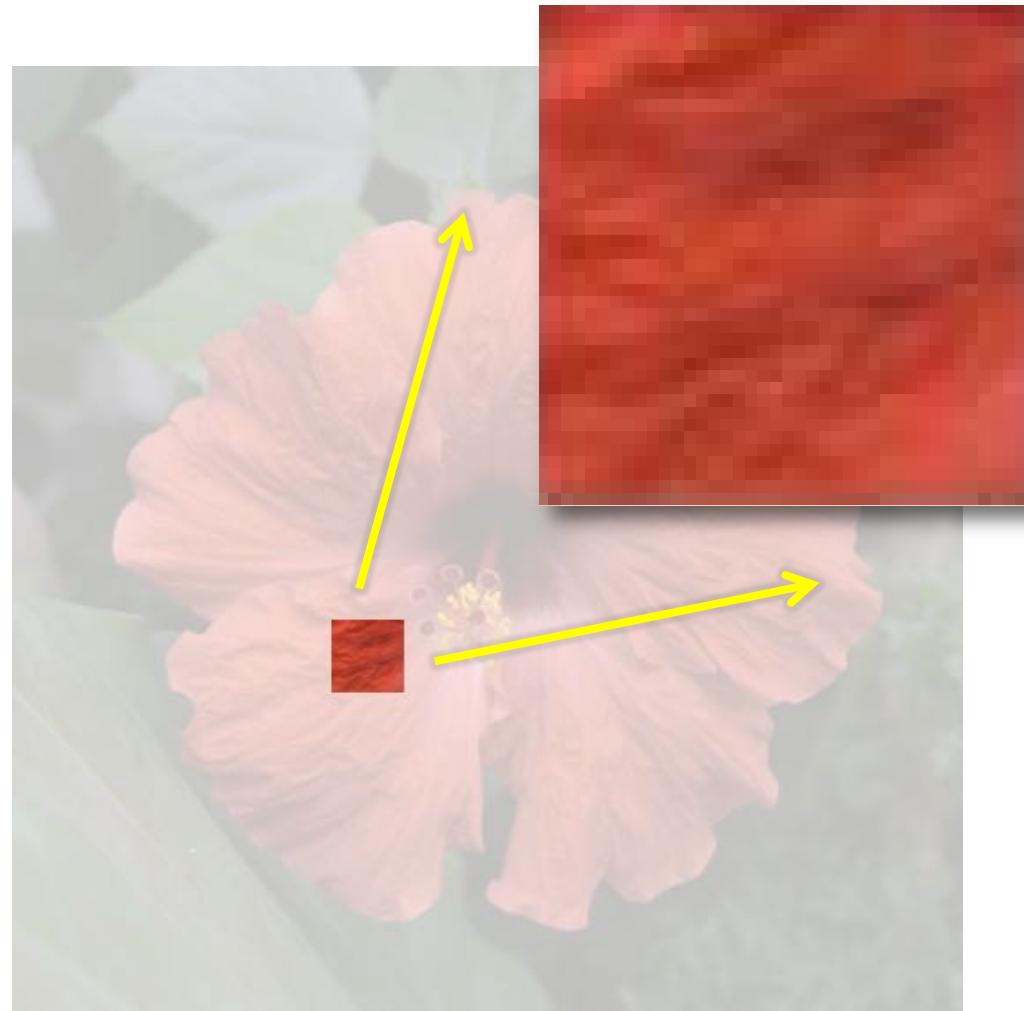
1. Generate an image that looks good for a small neighborhood
  - It may look bad elsewhere
2. “Combine data from all neighborhoods”

# Example: Local Contrast Increase



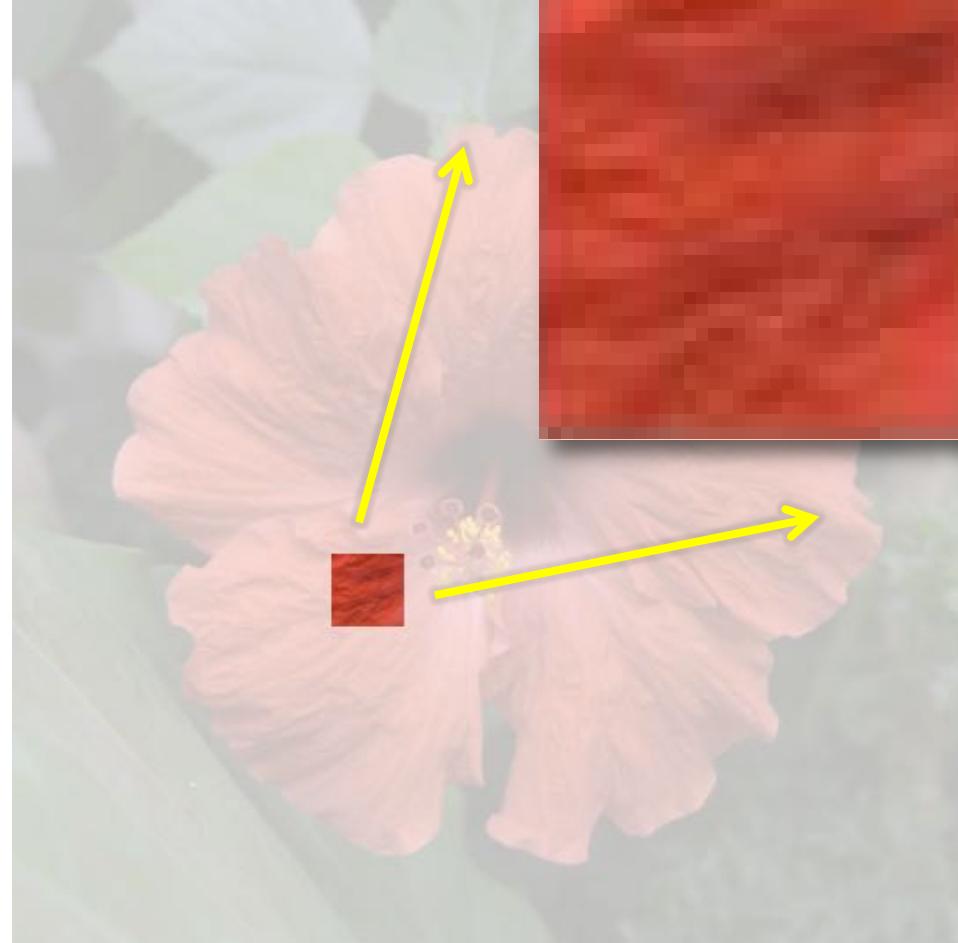
input

# Example: Local Contrast Increase

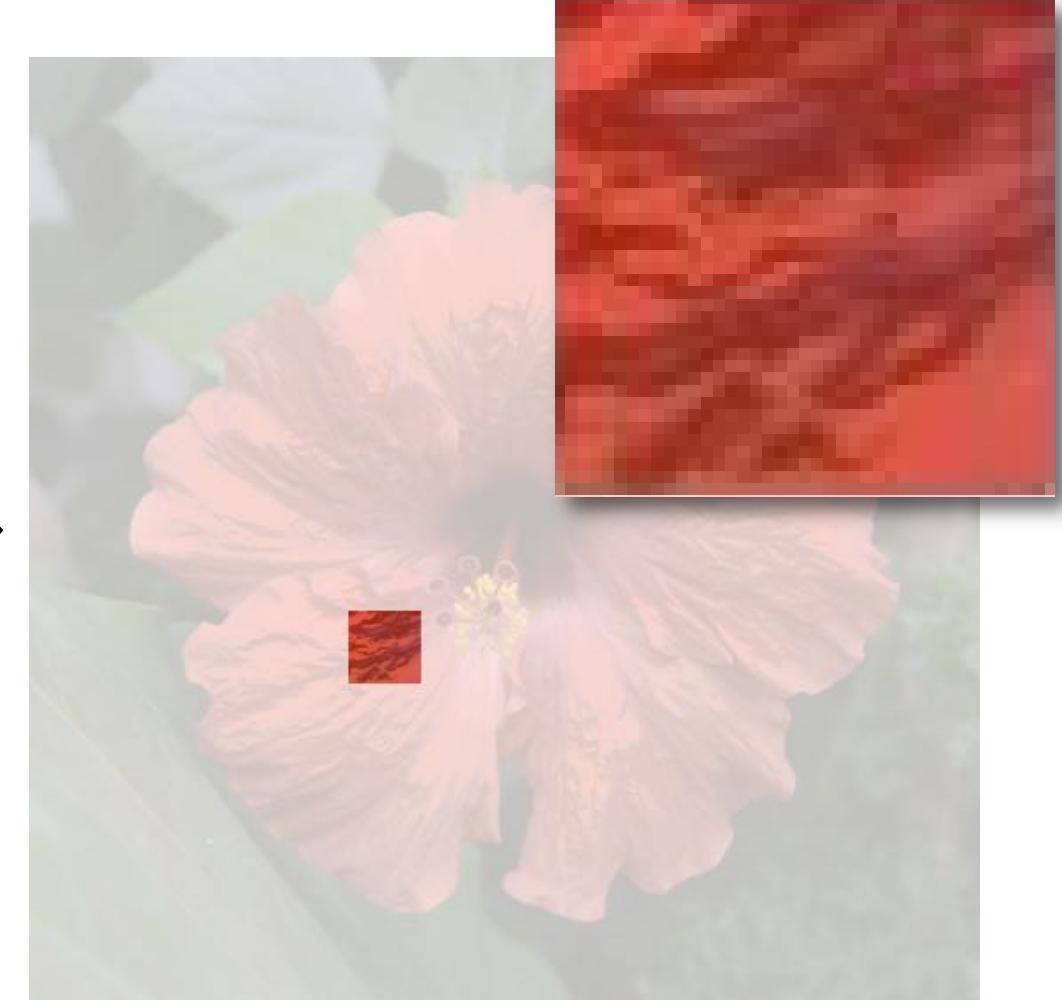


input

# Example: Local Contrast Increase

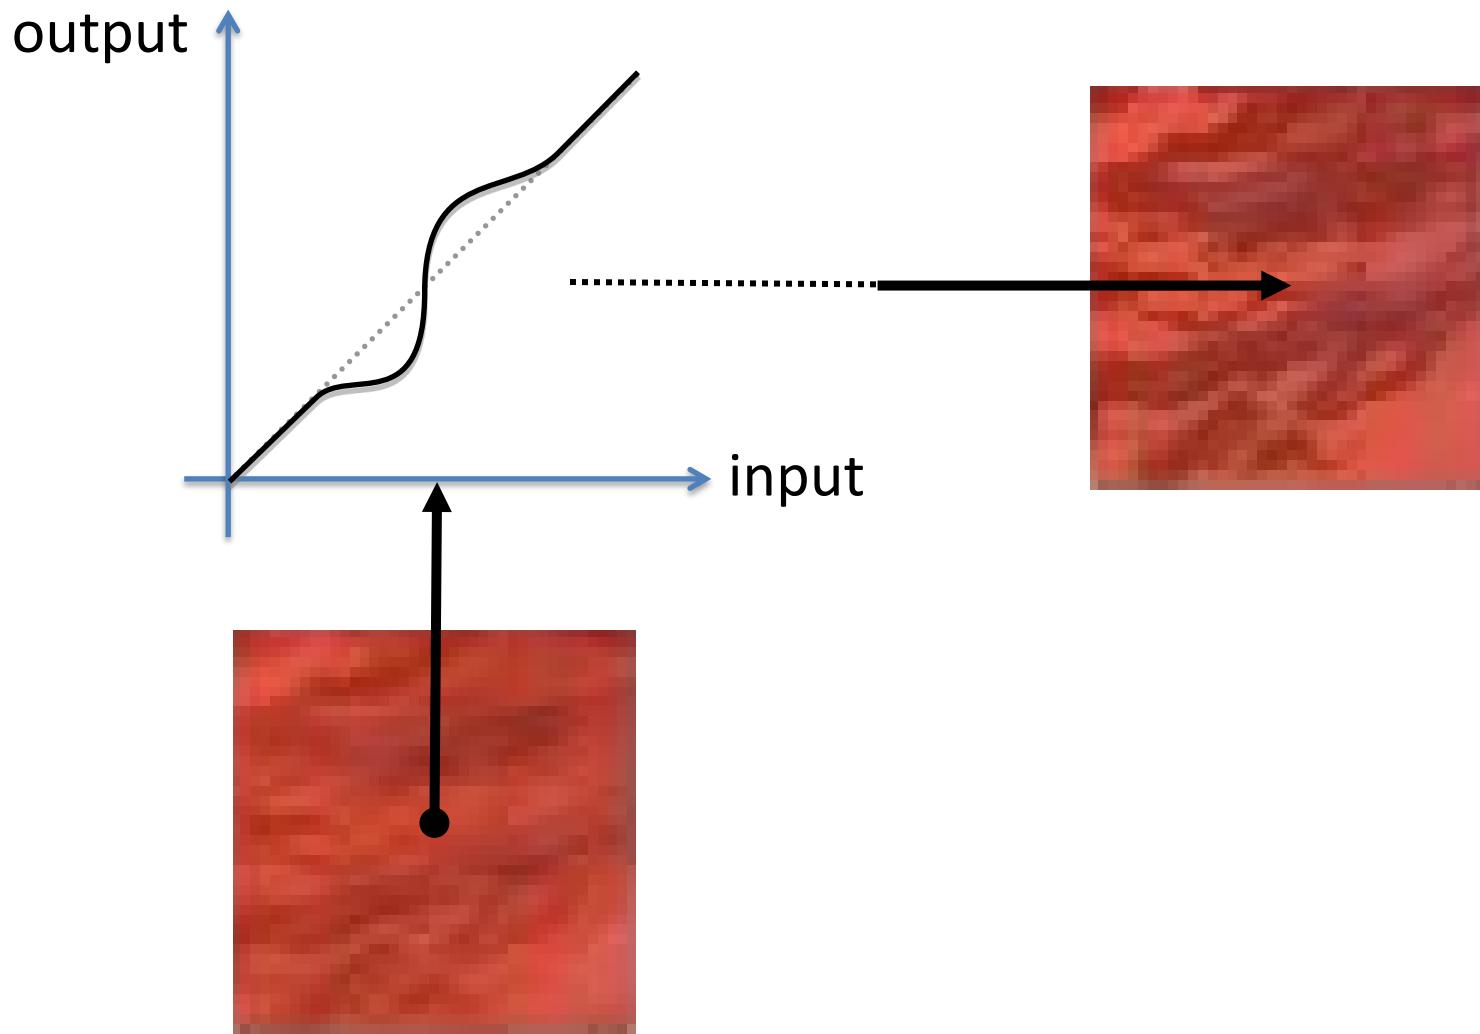


input



output with local contrast increased

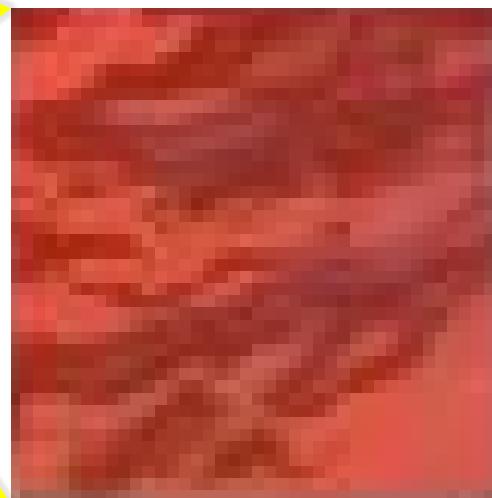
# Simple Local S-shaped Curve



# Only Local Result Matters



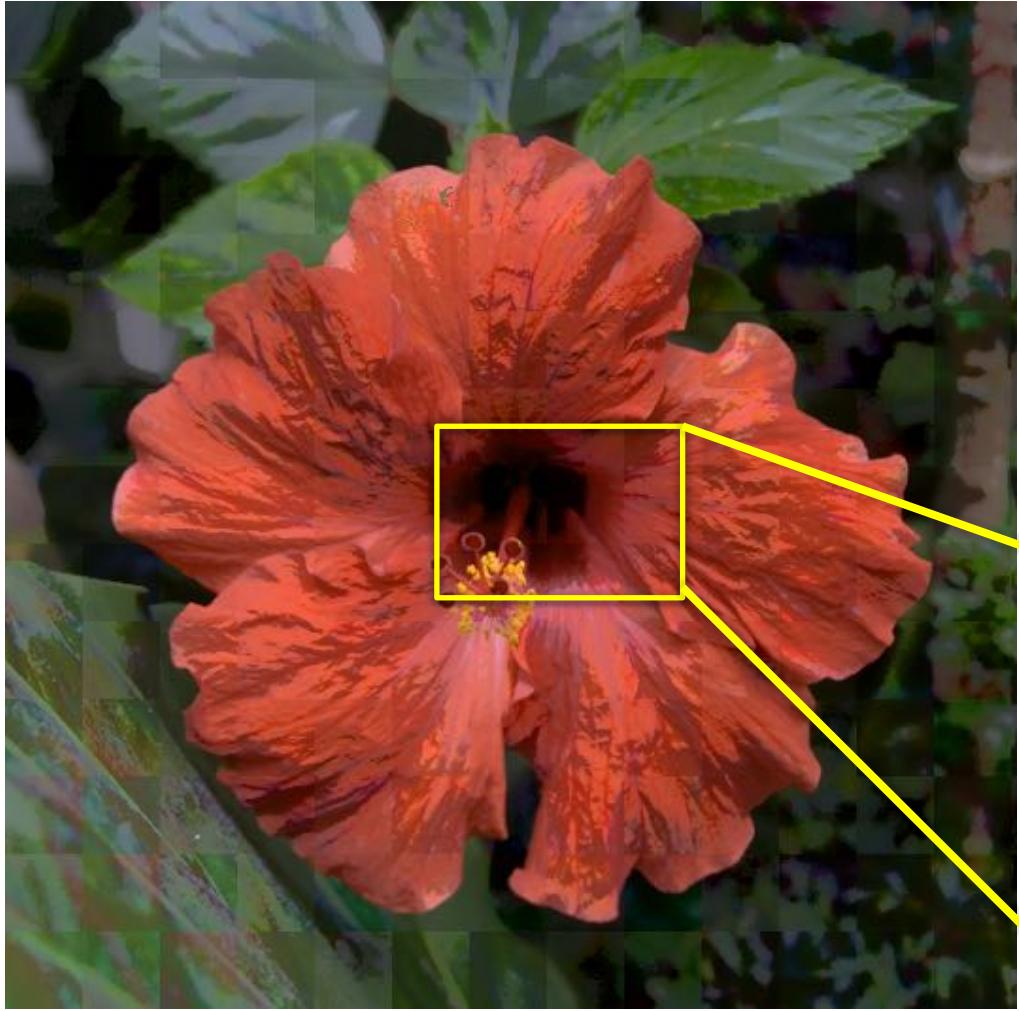
- Artifacts appear elsewhere
- Not a problem, we use only local data



output with local contrast increased

The processed image only needs to look good locally

# Naïve Stitching



- Paste local results side by side?
  - Visible seams, artifacts...
  - Possible heuristics: vary patch size, smooth seams...
  - But *ad hoc* and brittle



# Our Multi-scale Approach

- We build the Laplacian pyramid of the output coefficient by coefficient

For each coefficient

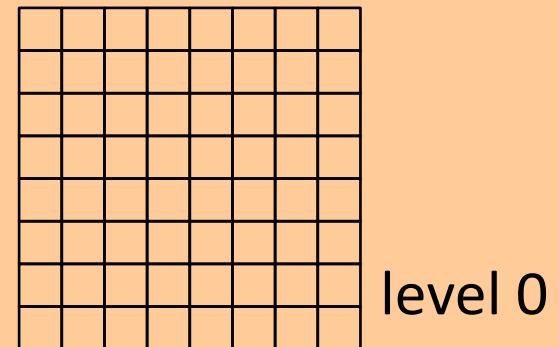
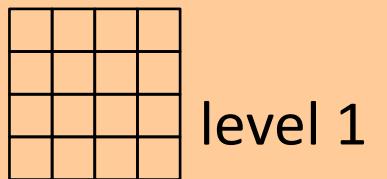
1. Generate “locally good” image
2. Compute Laplacian pyramid of that image
3. Copy coefficient to output pyramid

# Illustration

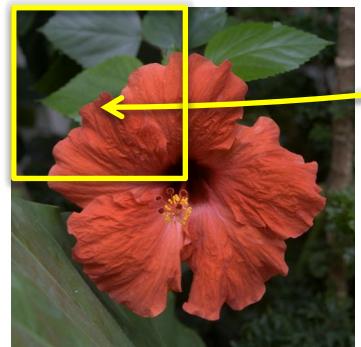


input image

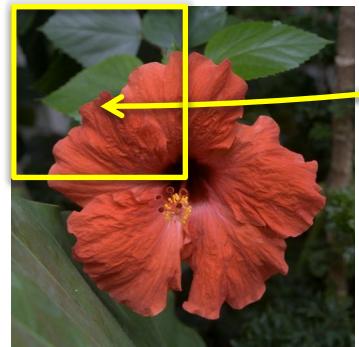
**Output  
Laplacian pyramid**



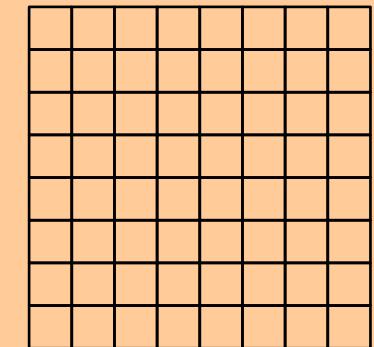
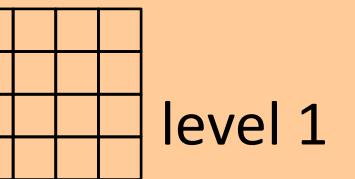
# Illustration



input image



**Output  
Laplacian pyramid**

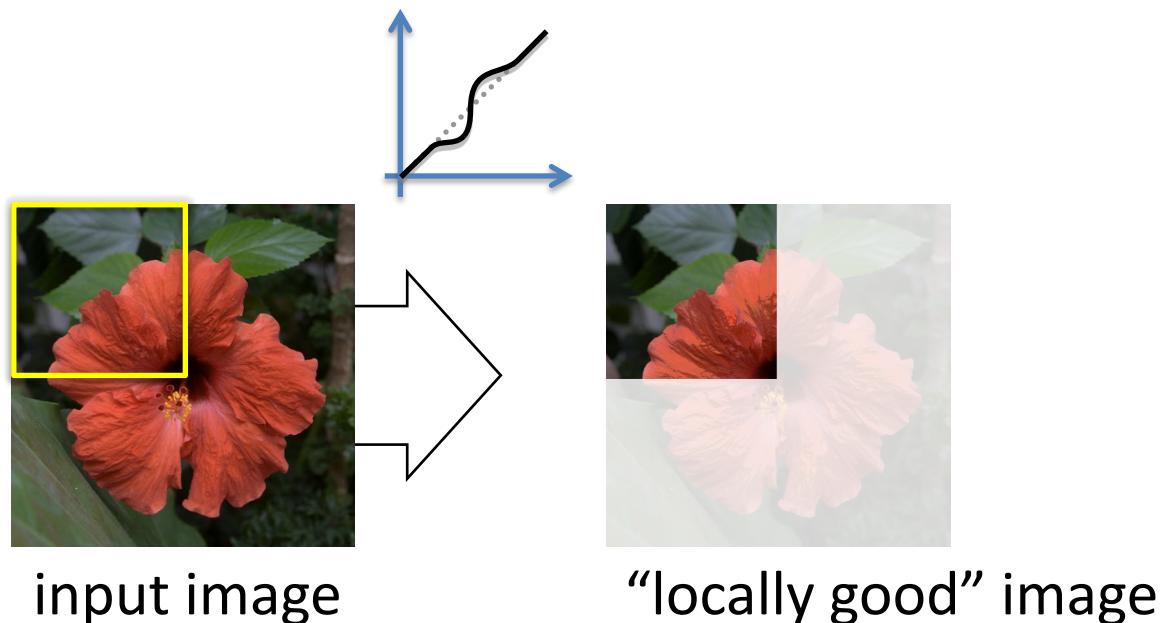


level 0

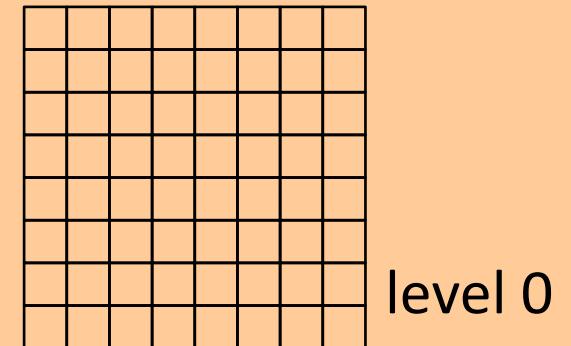
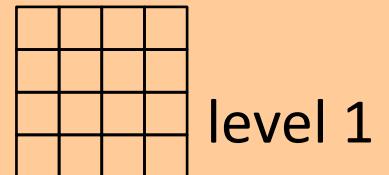
level 2

level 1

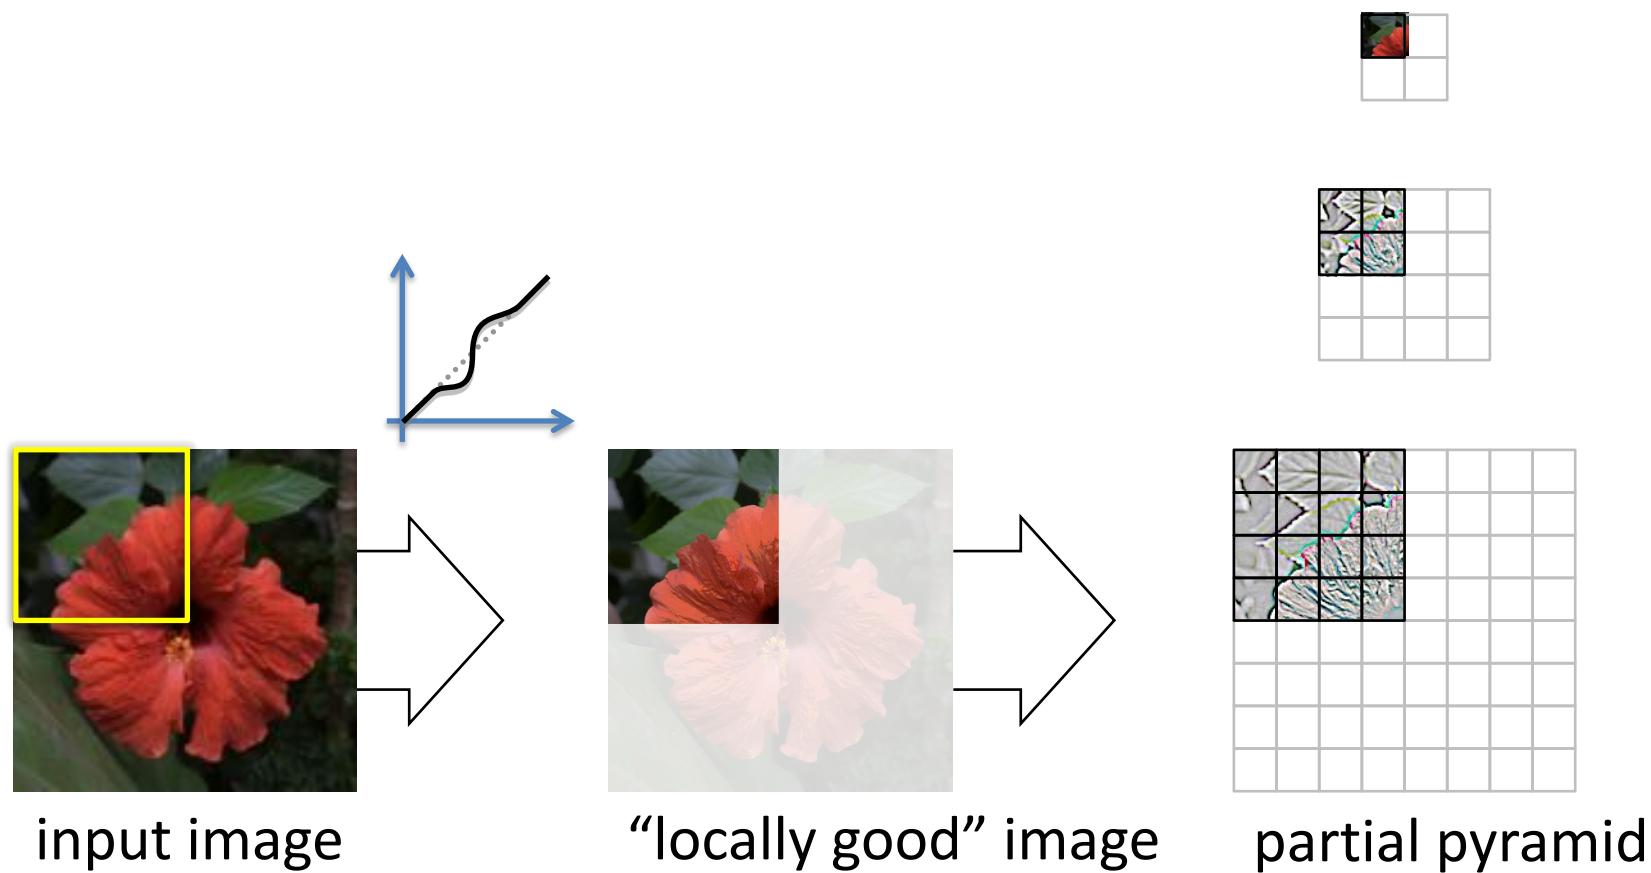
# Illustration



**Output  
Laplacian pyramid**



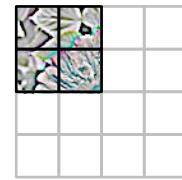
# Illustration



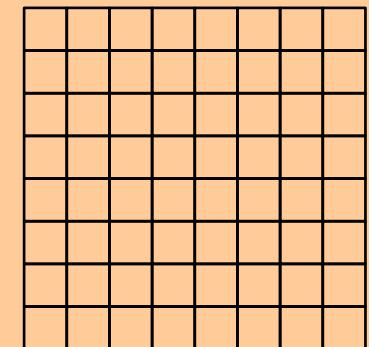
**Output**  
**Laplacian pyramid**



level 2

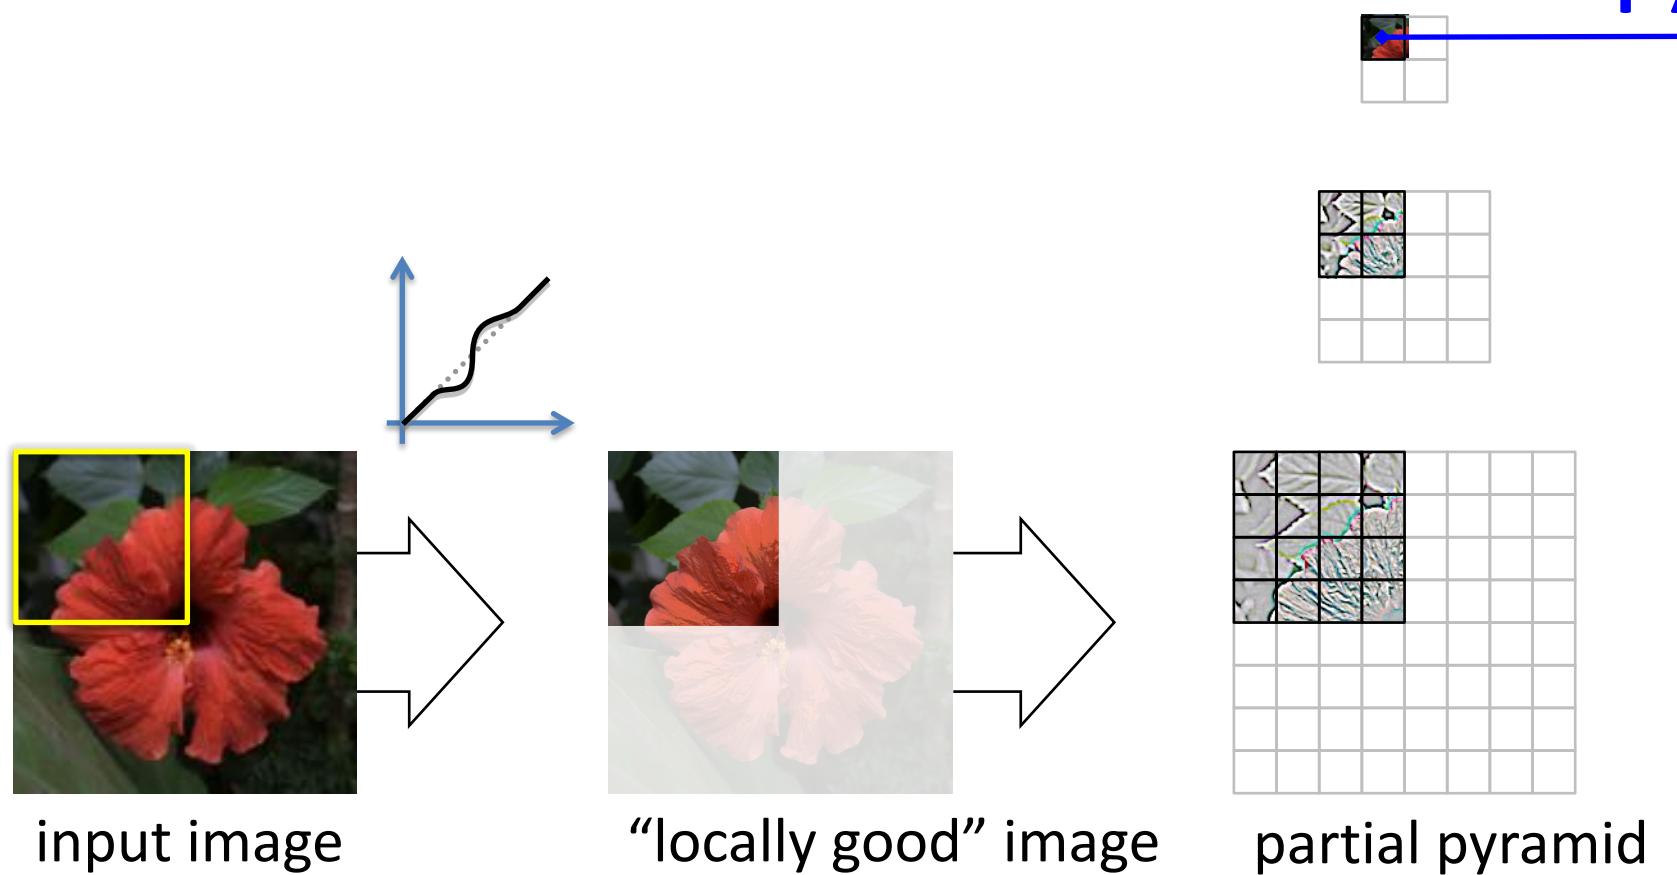


level 1



level 0

# Illustration

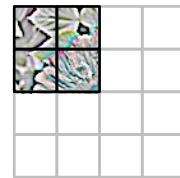


**Output  
Laplacian pyramid**

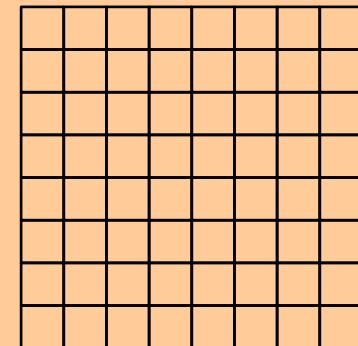
**copy**



level 2

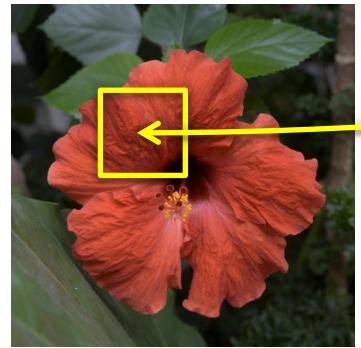


level 1

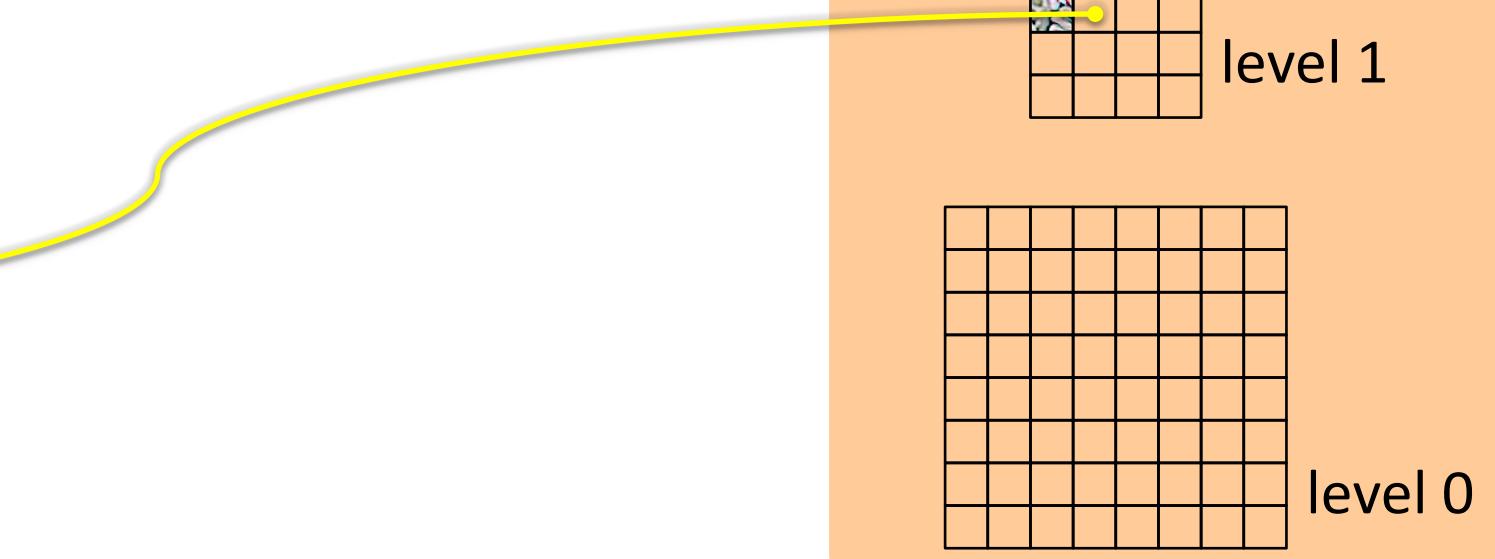


level 0

# Illustration

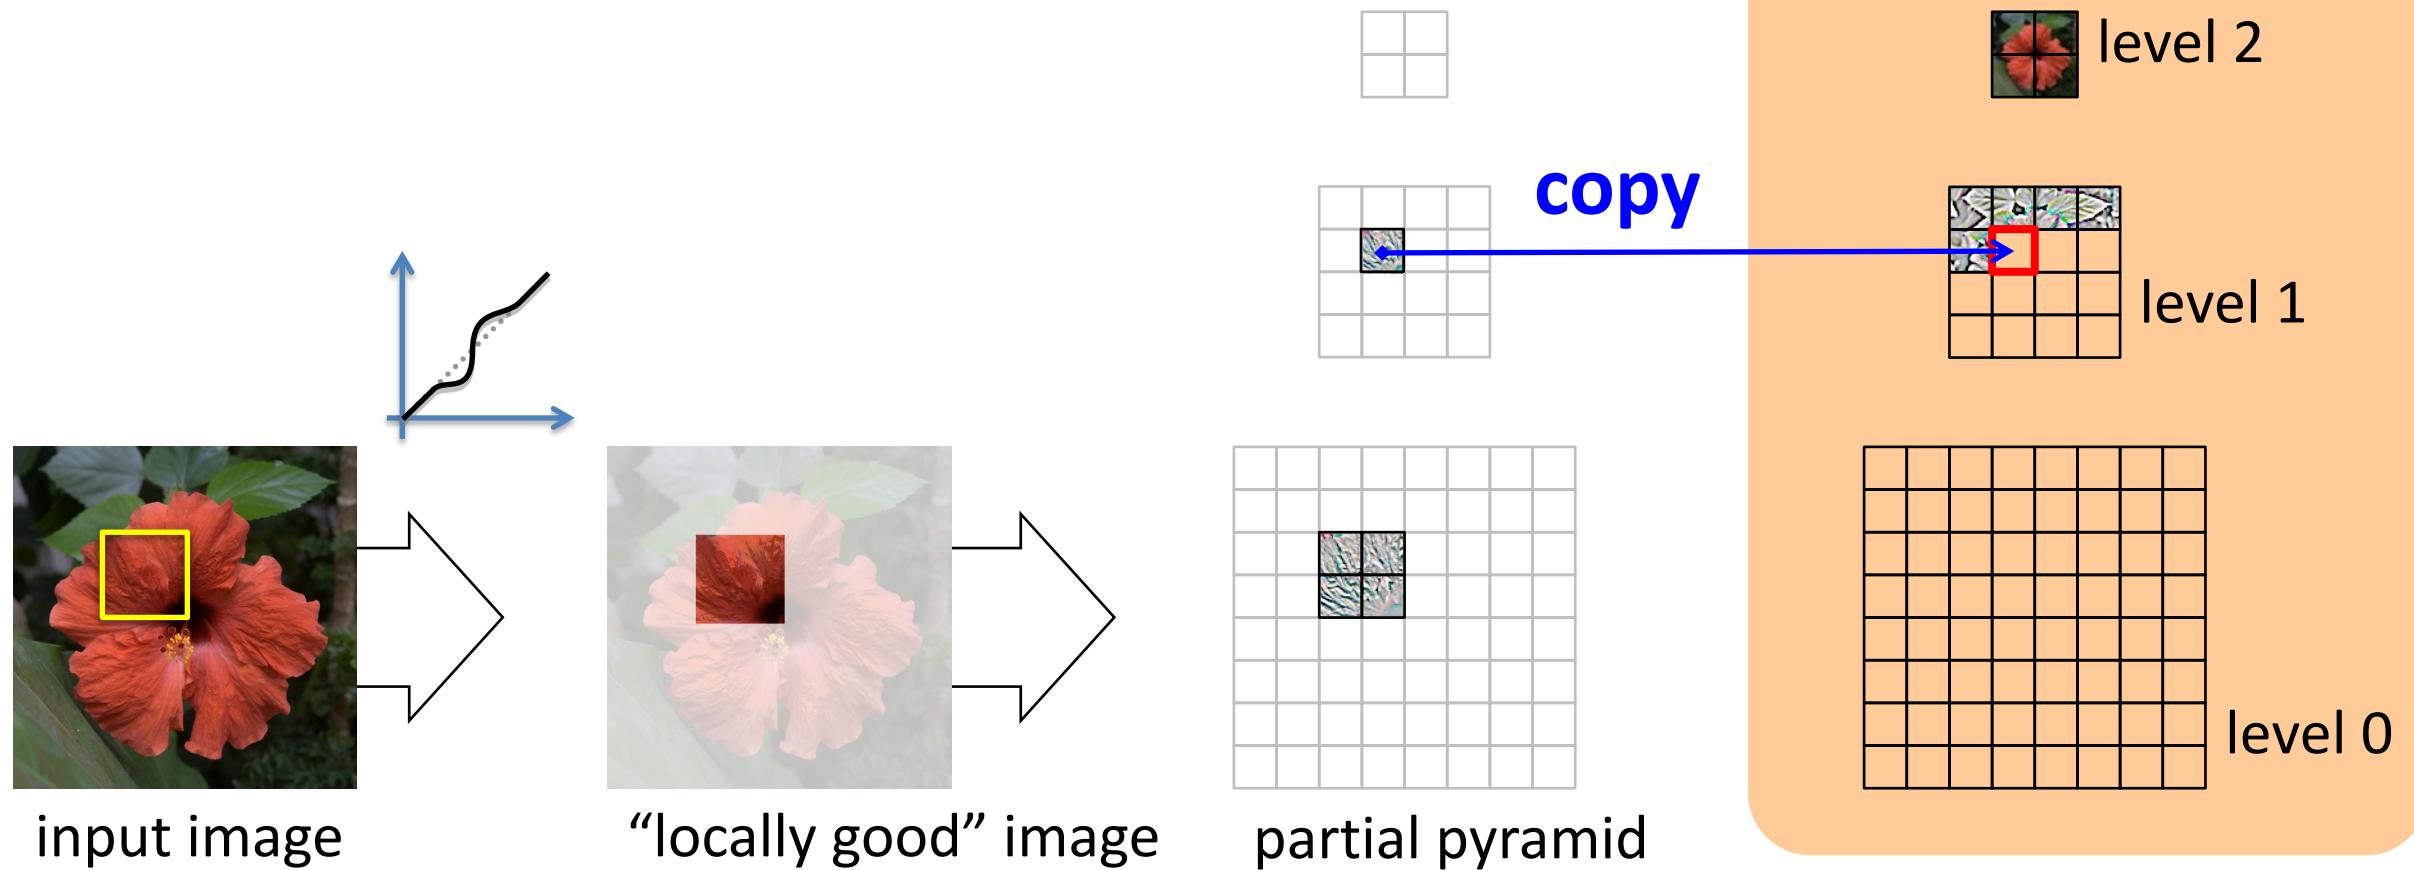


input image

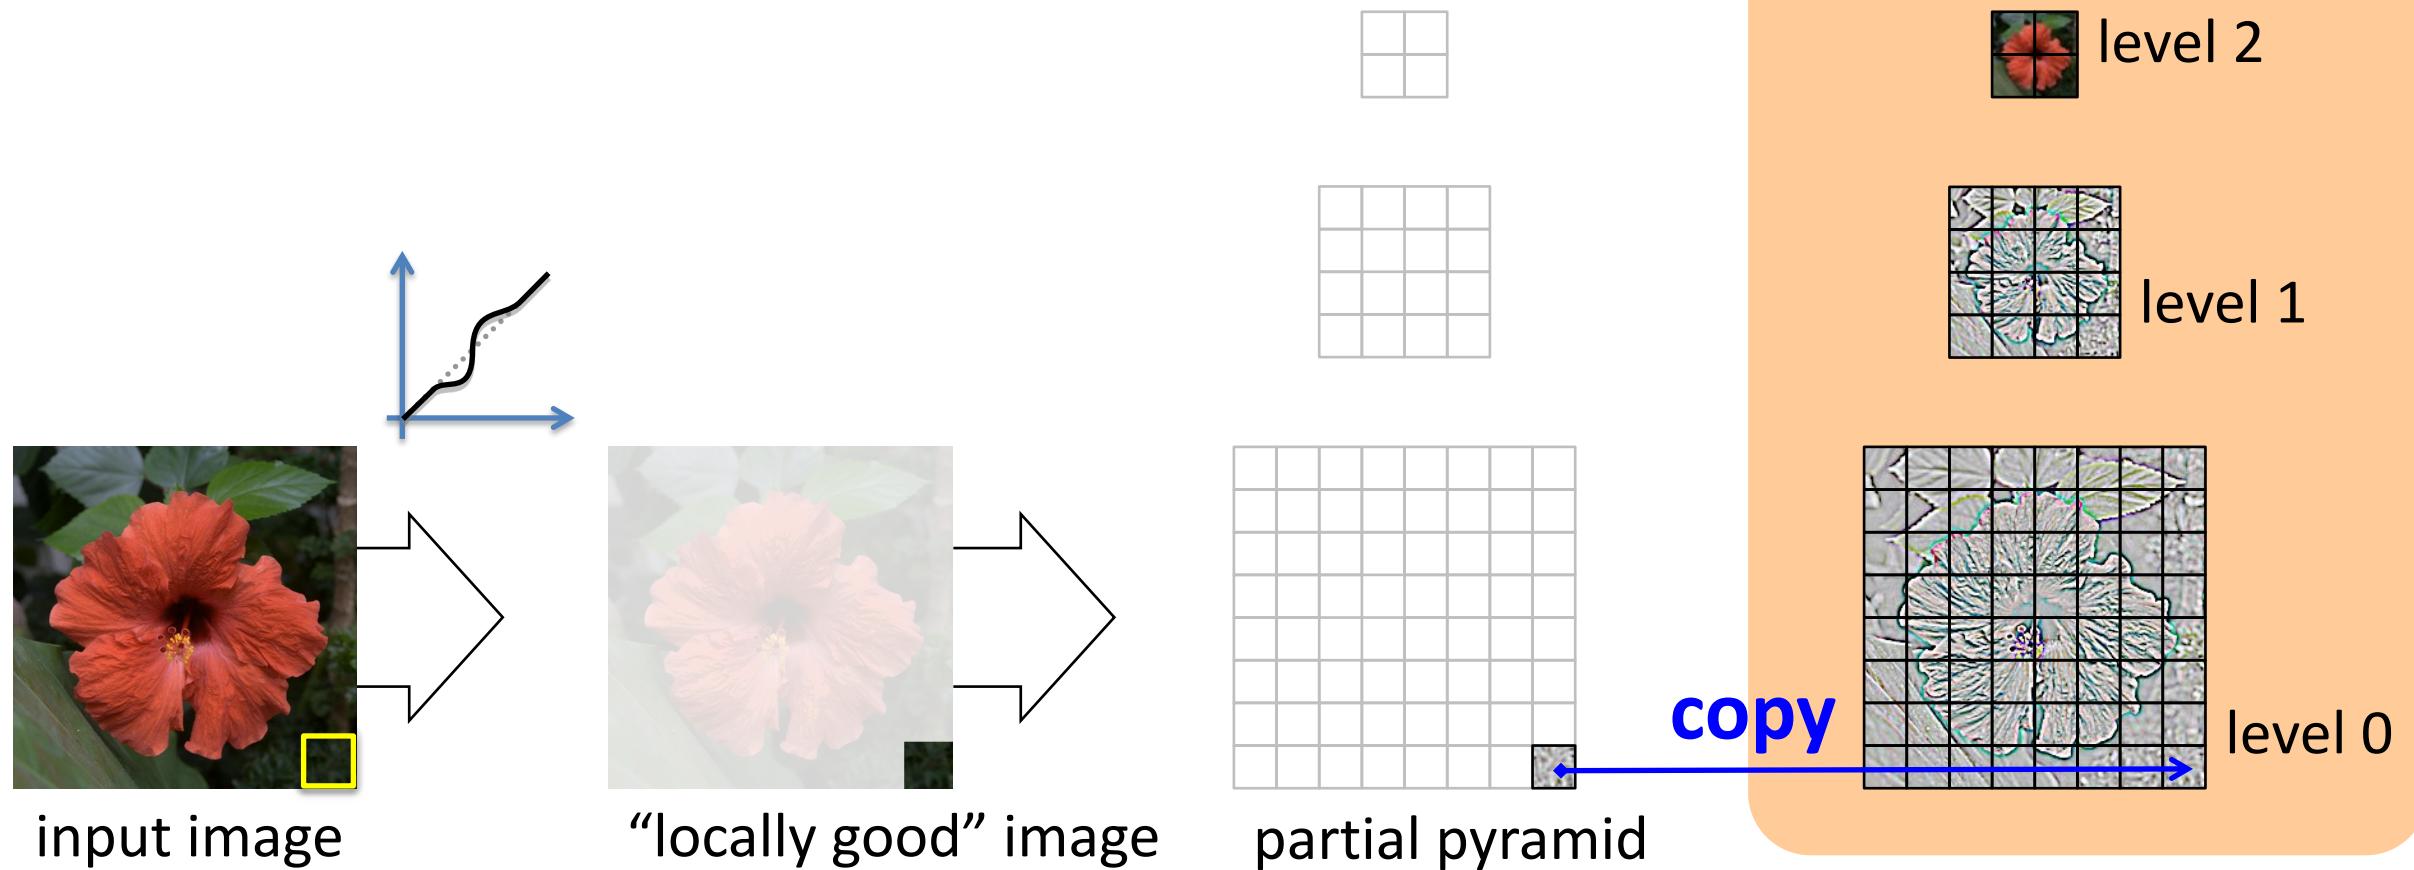


**Output  
Laplacian pyramid**

# Illustration

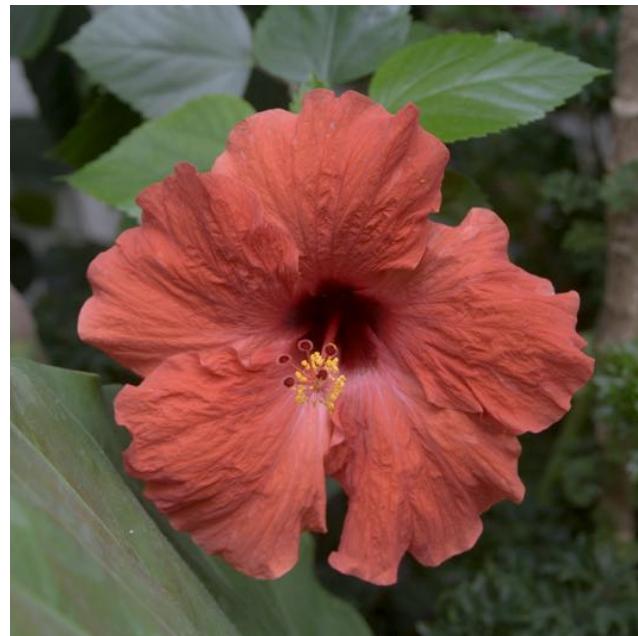
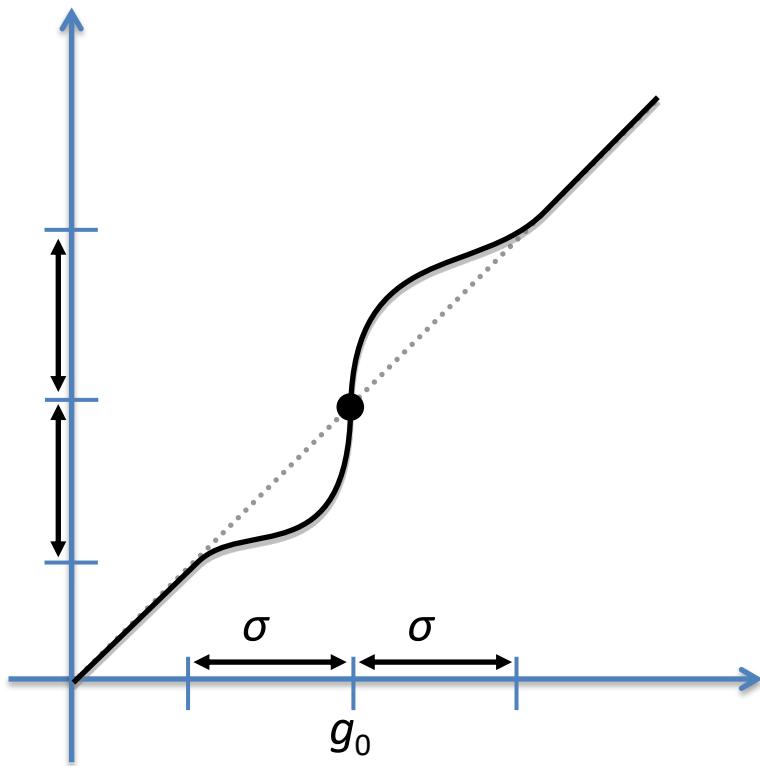


# Illustration



# Possible Nonlinearities

- Detail manipulation



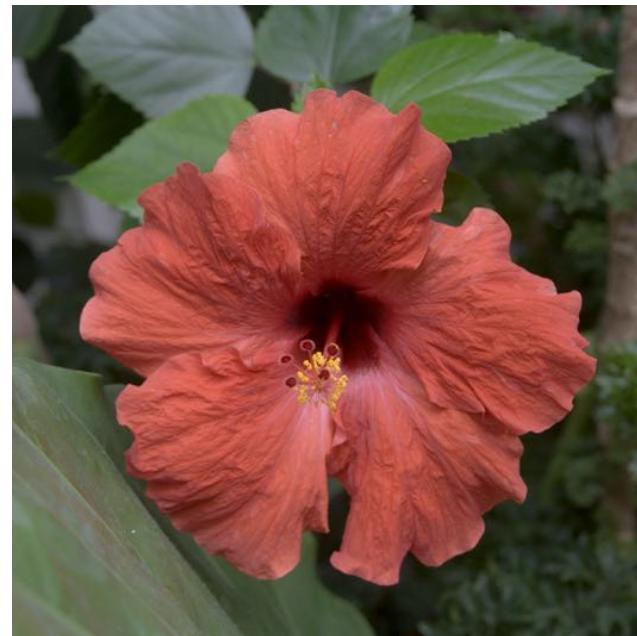
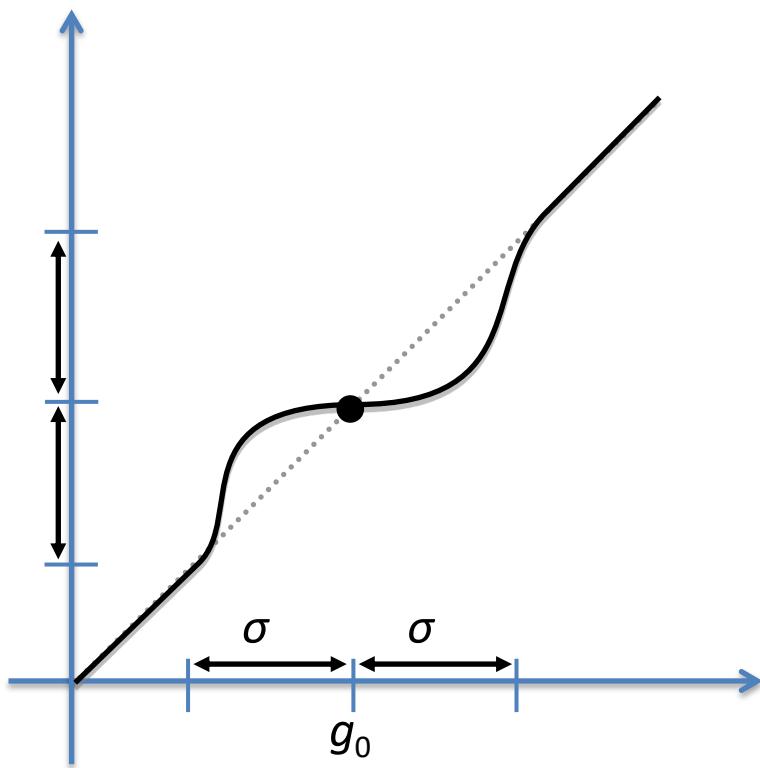
input image



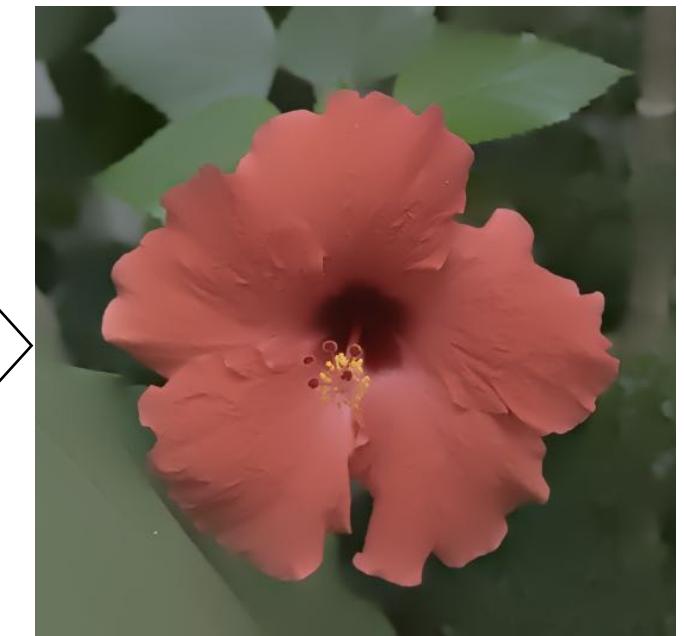
details enhanced

# Possible Nonlinearities

- Detail manipulation



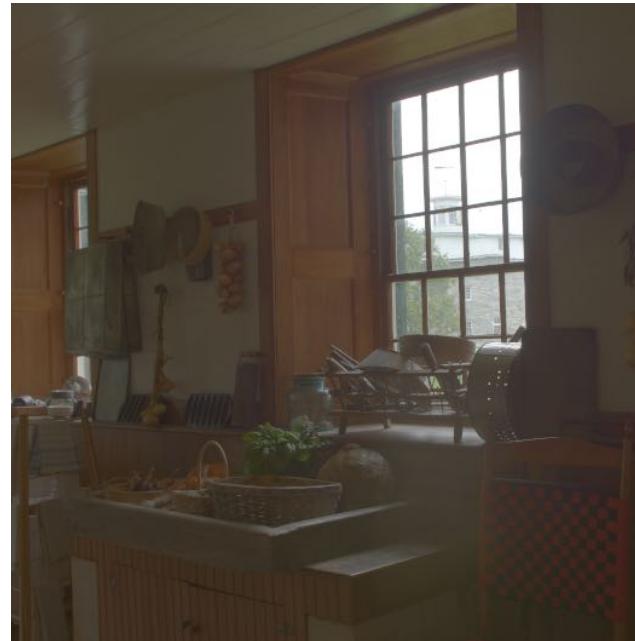
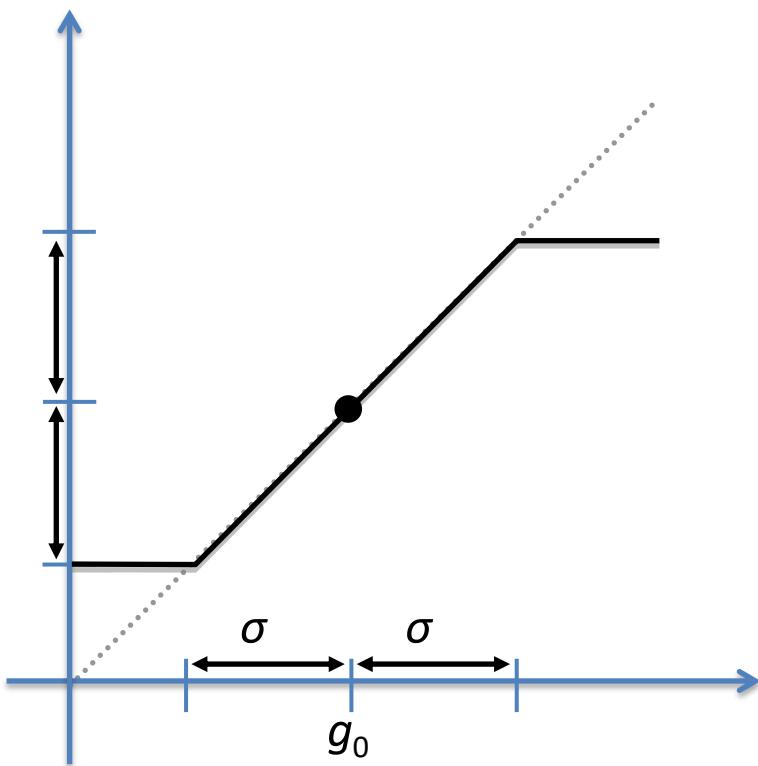
input image



details reduced

# Possible Nonlinearities

- Dynamic range manipulation



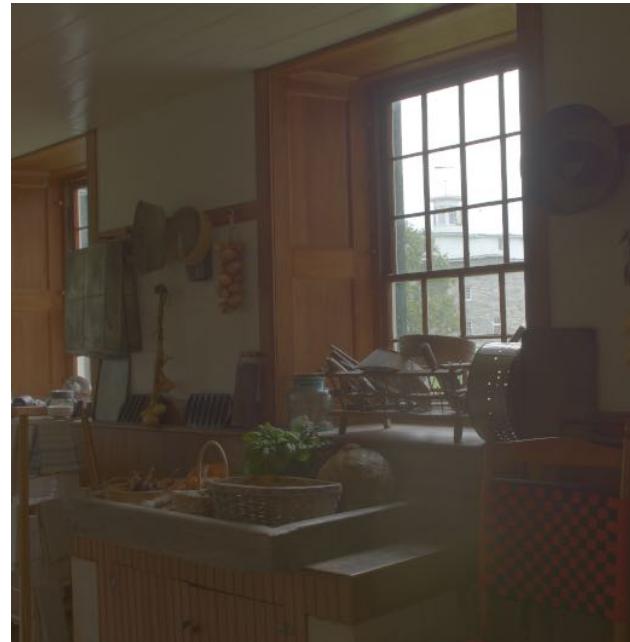
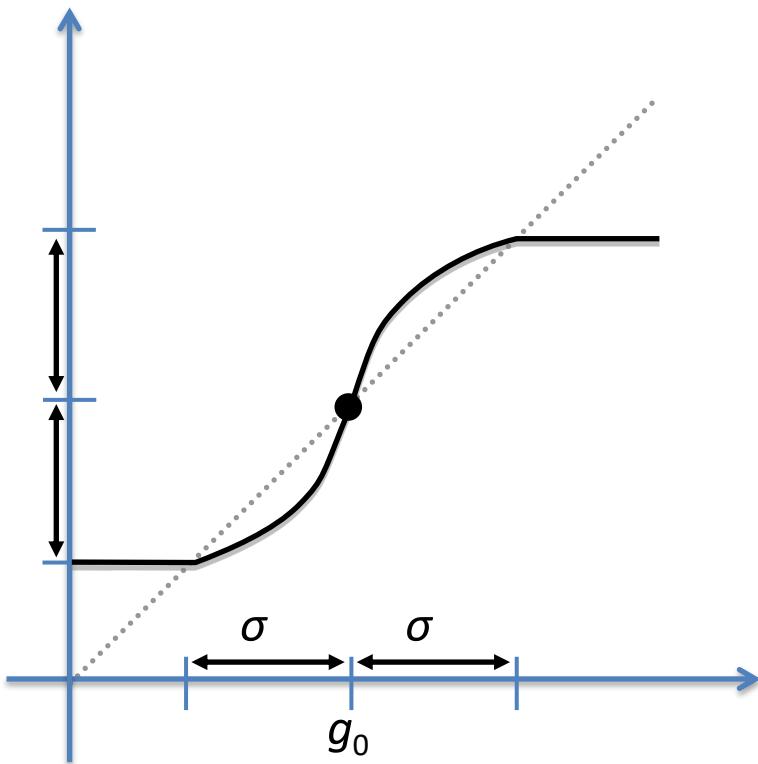
input image



tone mapped

# Possible Nonlinearities

- Can be combined, e.g. tone map + boost details



input image



tone mapped +  
details enhanced

# Photoshop Demo

# Back to Photo Style Transfer

*Transactions on Graphics 2014*



# Old Transfer Algorithm



# New Transfer Algorithm

# Close-up



old algorithm



new algorithm  
*a lot more details  
higher quality*

Takeaway message #1

**Photo style transfer works.**

Takeaway message #2

**Low-level aspects matter.**

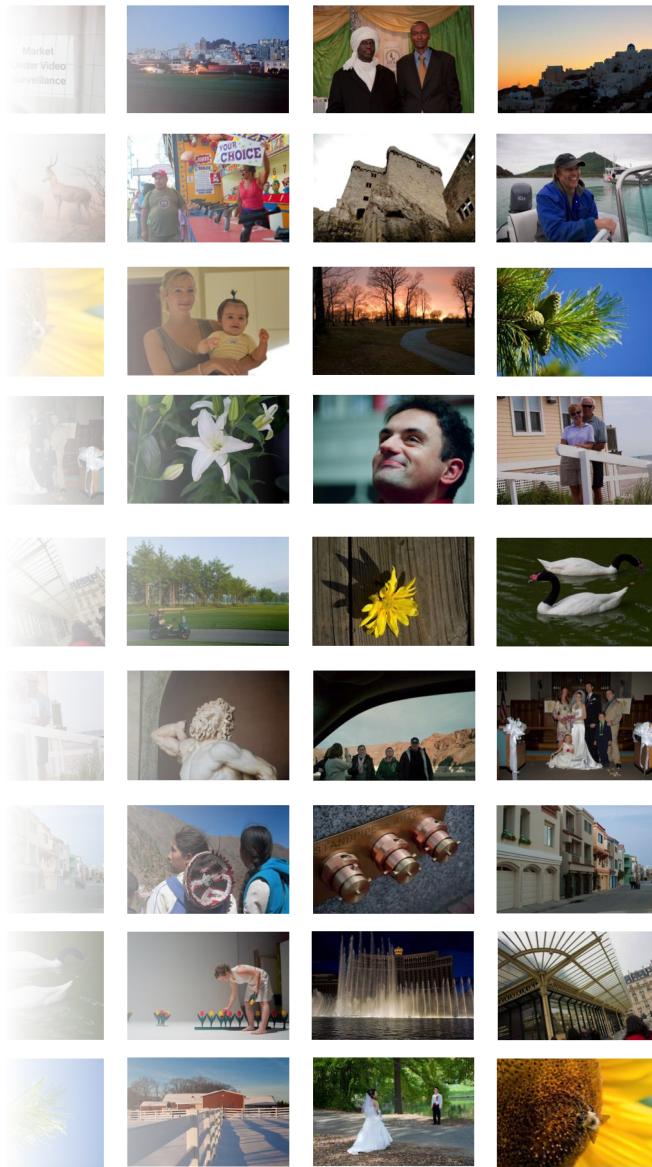
# Learning from a Dataset

*CVPR 2011*

5000 reference photos  
adjusted by a pro



input



5000 reference photos  
adjusted by a pro



input



5000 reference photos  
adjusted by a pro



our result



# Result



# Result



# Product Impact

- Transferred into Photoshop CS6 as “Auto” in Brightness & Contrast, Levels, Curves, and Auto Tone. 
- Can also learn from users on the fly, transferred as Auto Smart Tone into Photoshop Elements 12. 

# Advantages of Gaussian Processes

- Can learn from a few hundred examples “only”
  - Hours of work by a professional → quality training data
- Revert to the mean when “unsure”
  - Does not ruin the image
- Easy to debug by looking at how training data are used

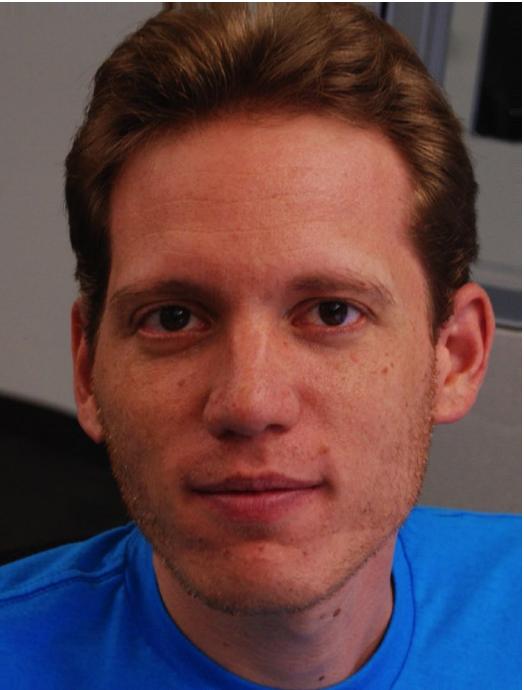
**Back to Transferring from an Example**

# Previous Techniques Fail on Portraits

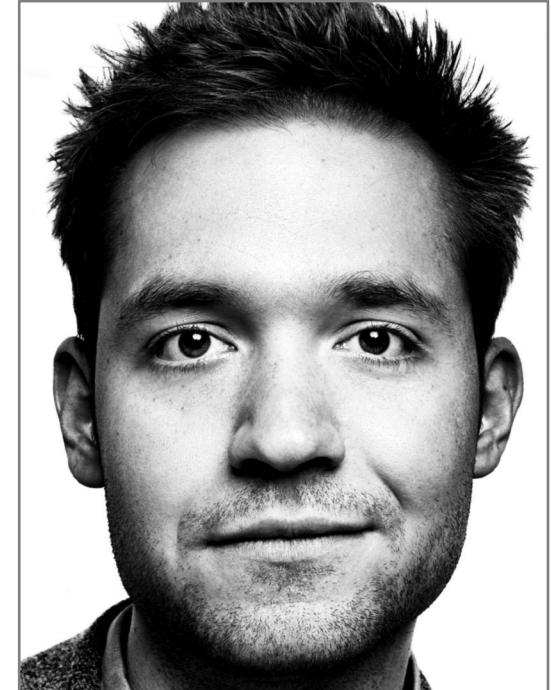
- Two-scale model not sufficient for skin texture
- Image-level statistics too coarse for portraits



Make this portrait



look like this one



# A Local and Multi-scale Model



Input



Target

# A Local and Multi-scale Model

1. Construct Laplacian stacks for the input and the example



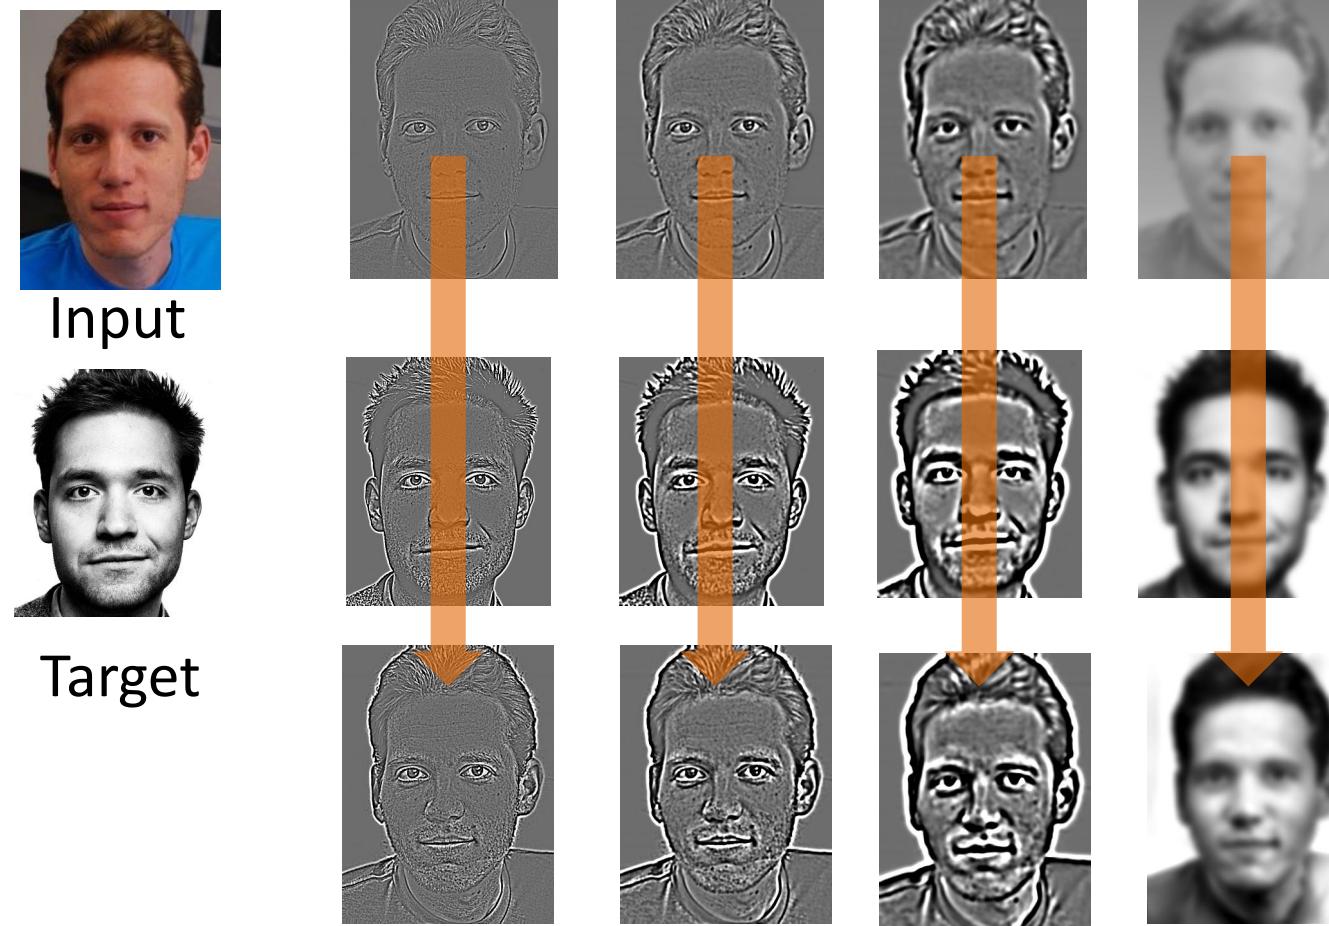
Input



Target

# A Local and Multi-scale Model

1. Construct Laplacian stacks for the input and the example



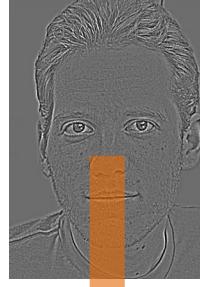
2. Local match  
at each scale

# A Local and Multi-scale Model

1. Construct Laplacian stacks for the input and the example



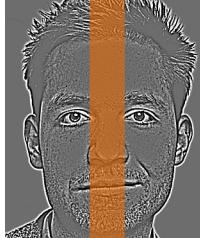
Input



2. Local match  
at each scale



Target

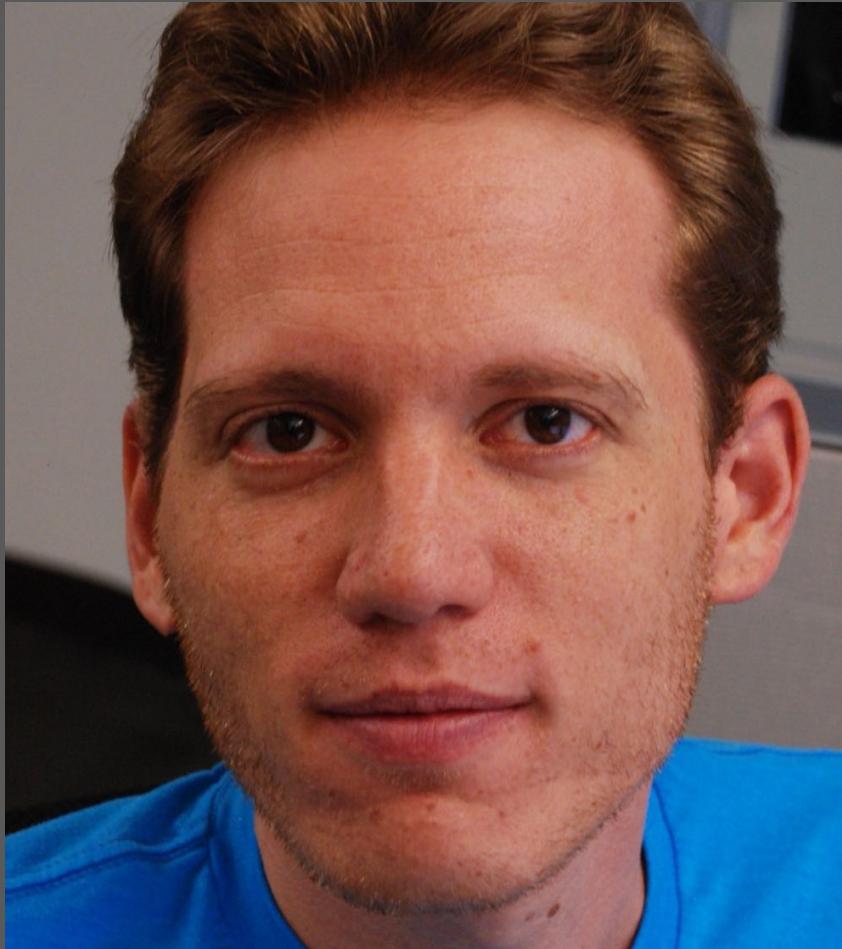


3. Collapse the matched stacks to create the output of this step



Output

# Result (Plato)



input

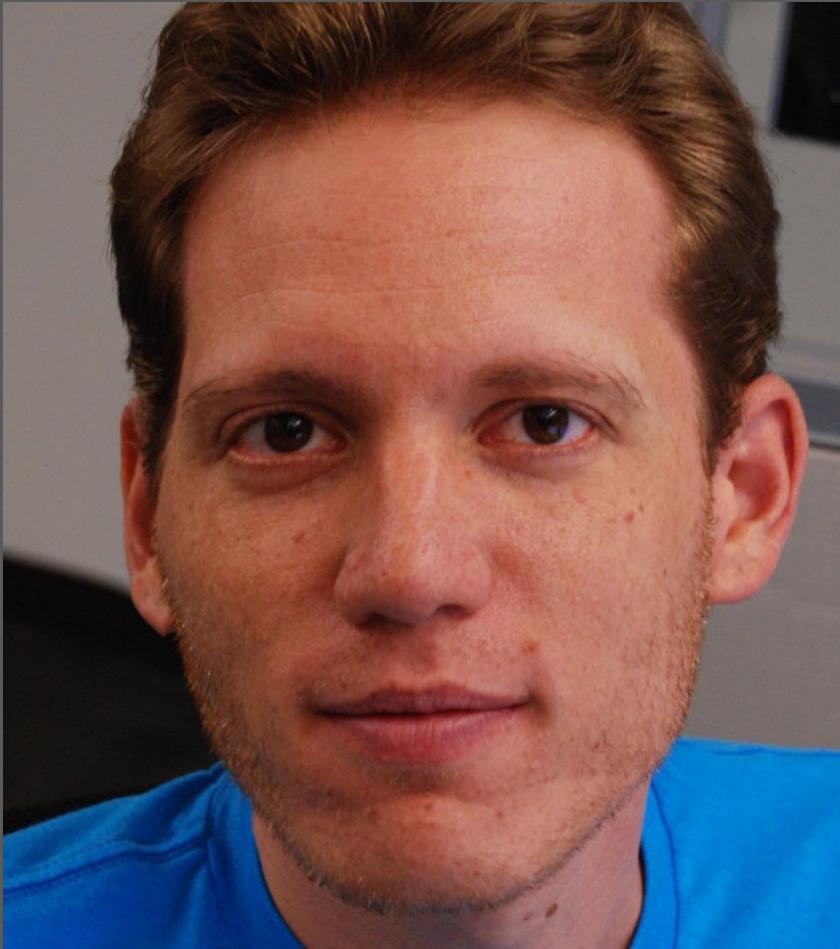


output



example

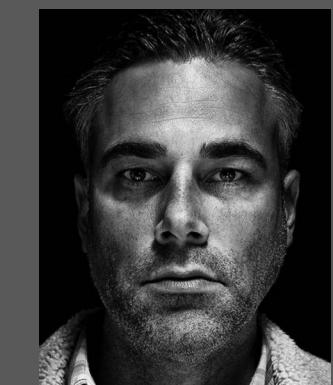
# Result (Kelly Castro)



input

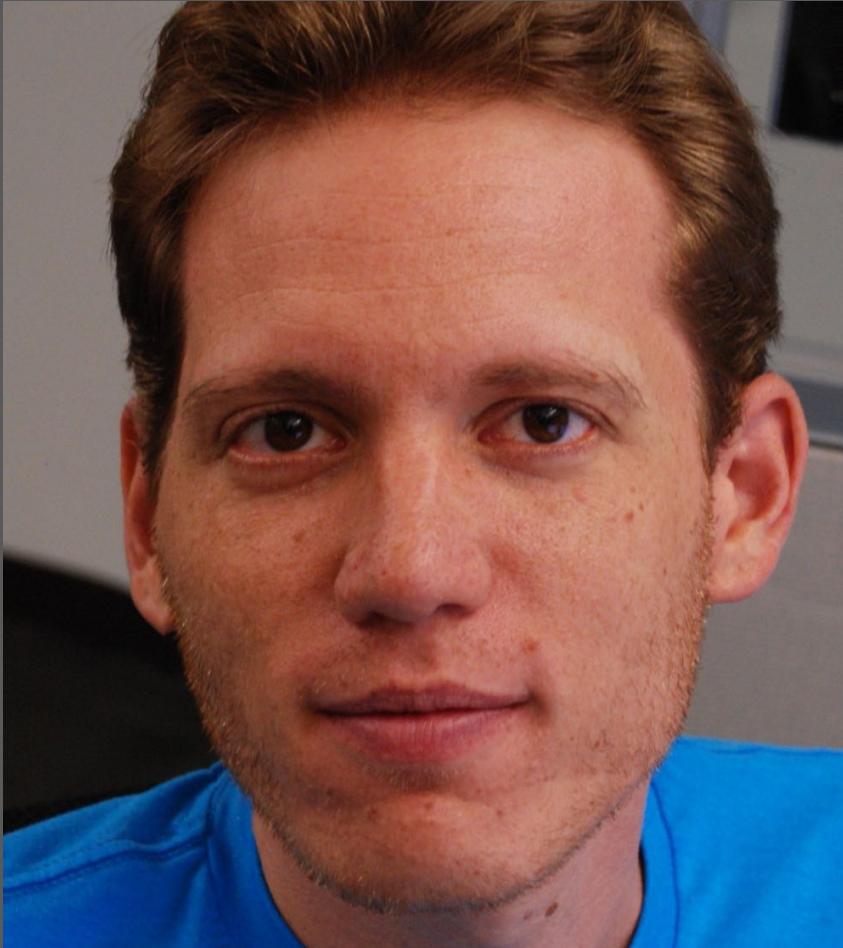


output



example

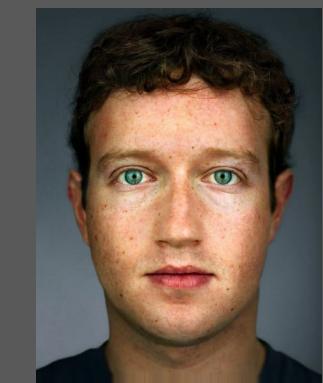
# Result (Martin Schoeller)



input



output



example



Input sequence with extreme facial expressions



Our style transfer result using the example in the gray box

Transferring Other Properties

# Time-of-Day Hallucination

*SIGGRAPH Asia 2013*

# Time-lapse Videos as Examples

- A database of 400+ time-lapse videos
- Shows how scenes change during the day





Input photo



Our result



Input



Our result (golden hour)





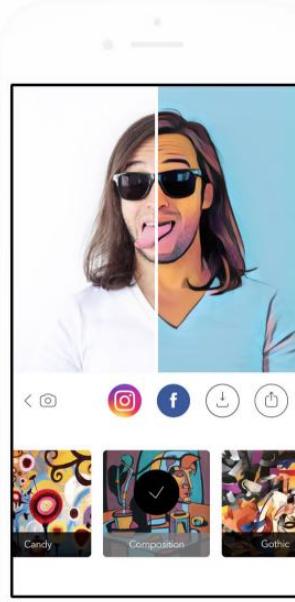
# **Deep Photo Style Transfer**

*A First Step Toward a Unified Approach*

*To appear at CVPR 2017*

# Motivation

- A Neural Algorithm of Artistic Style  
[Gatys et al. 2016]
- Several apps and websites
  - Prisma
  - DeepArt.io



The neural style algorithm works well for paintings.

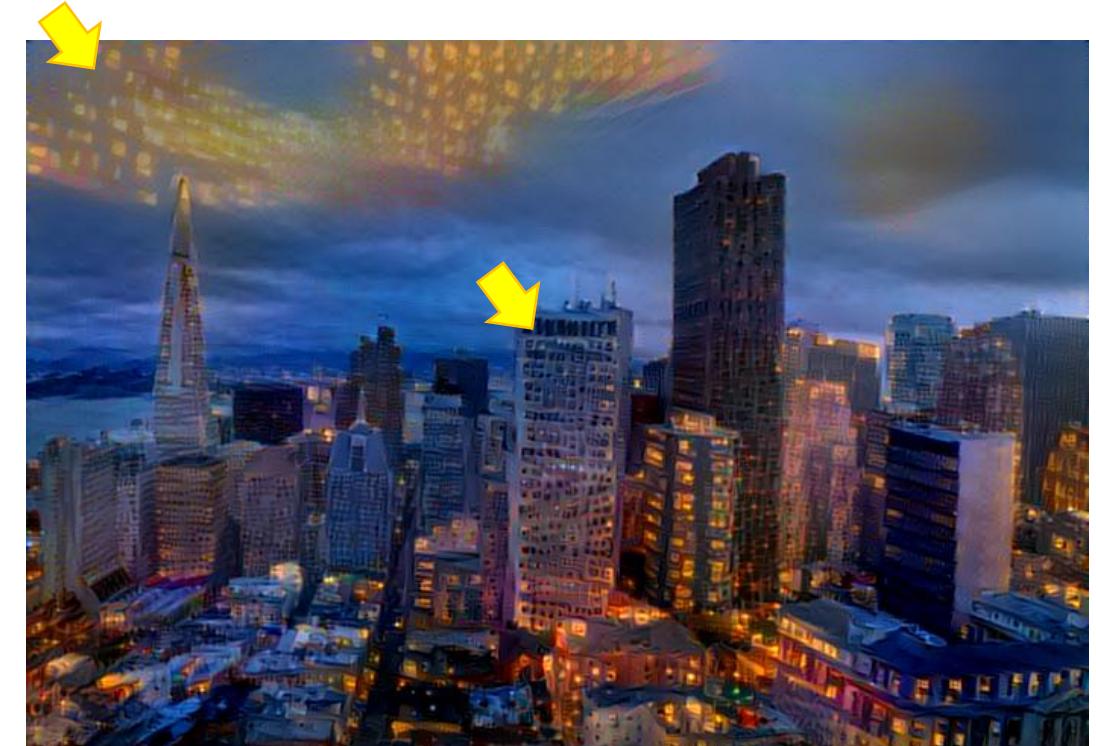
What about photos?

So we tried...



Input

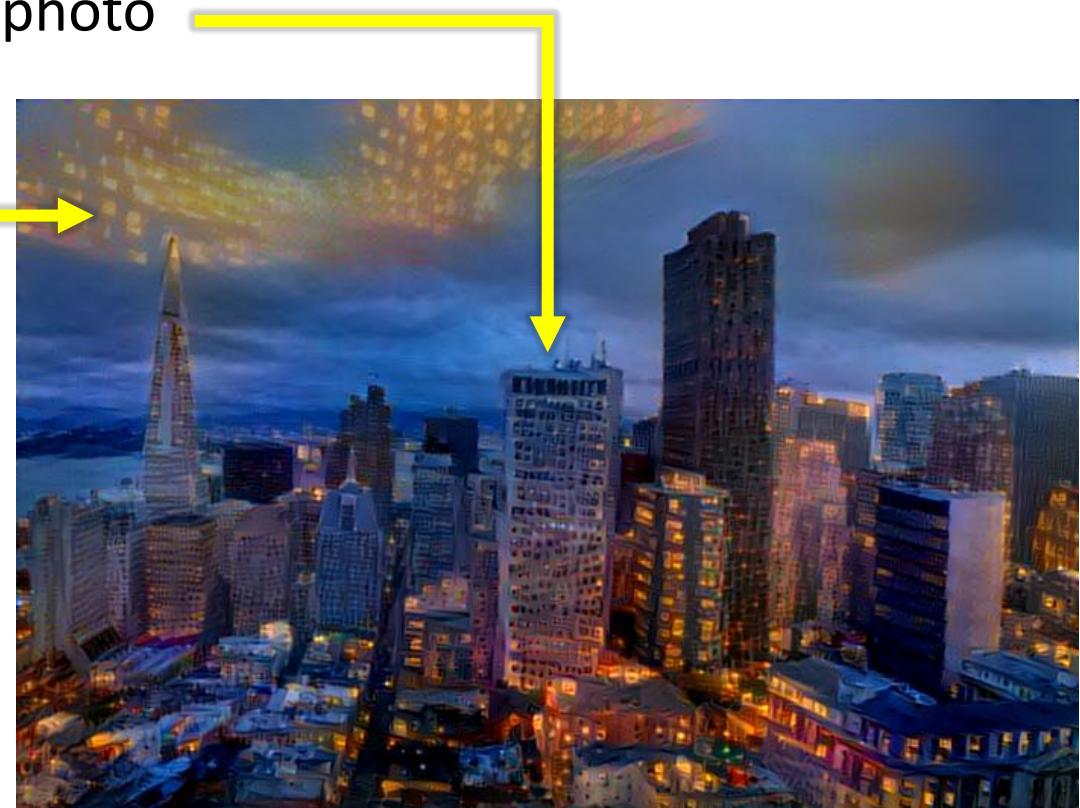
Model



Output of Neural Style Transfer

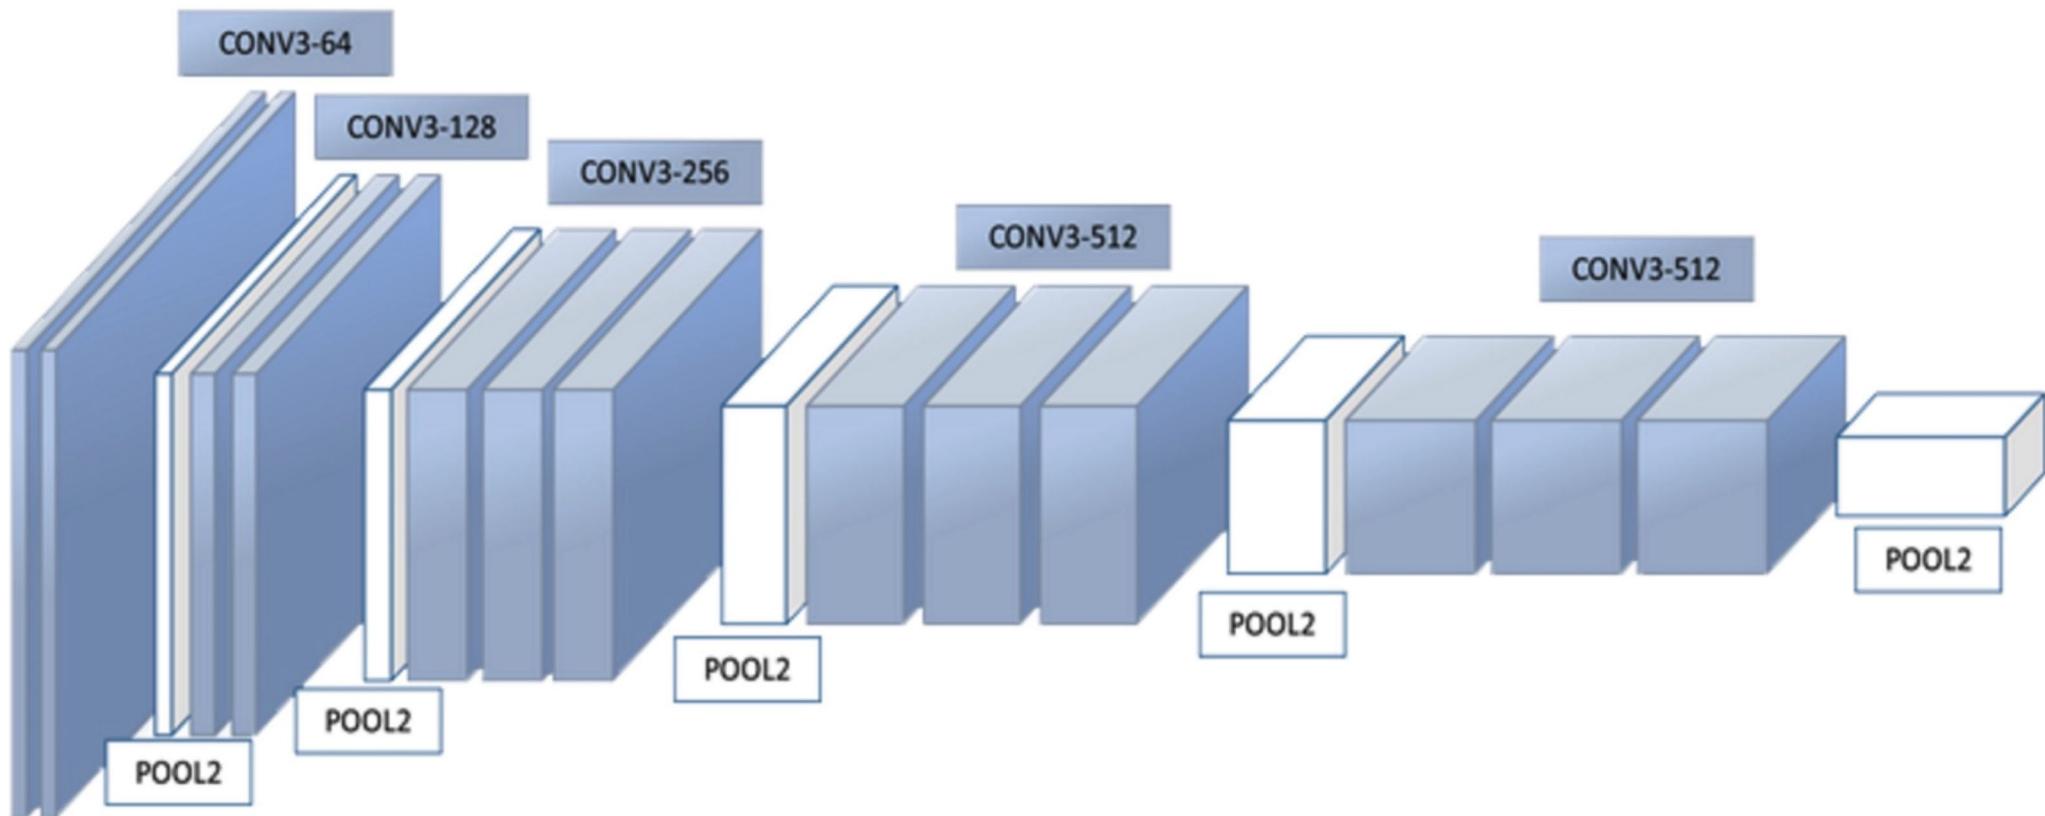
# Two Problems to Solve

- Undesired distortion
  - Result still looks like a painting, not a photo
- Semantic mismatch
  - E.g., ground texture can appear on the sky



# Background on Neural Style Transfer

1. Decompose the input and model images using a neural network



# Background on Neural Style Transfer

1. Decompose the input and model images using a neural network
2. Impose some of the statistics of the model image

# Background on Neural Style Transfer

1. Decompose the input and model images using a neural network
2. Impose some of the statistics of the model image
3. Reconstruct the output image

# Background on Neural Style Transfer

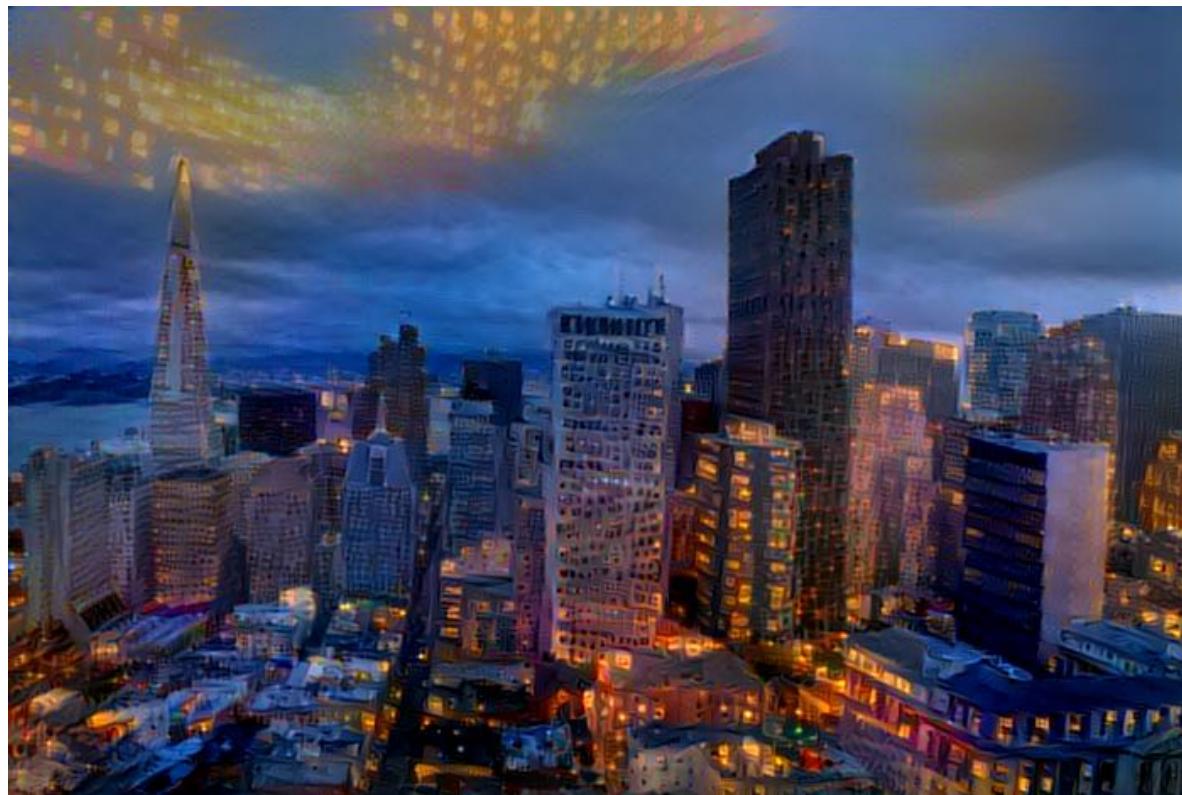
1. Decompose the input photo using a neural network
2. Impose some of the statistics of the model image
3. Reconstruct the output image

Preventing distortion

# Preventing Distortion

- We tried many options
  - Constraining the gradients
  - Multi-scale constraints akin to Portrait Style Transfer
  - Band-limiting the transformation
- All helped but none was 100% successful
- What worked was forcing the color transformation to be locally affine  
[Levin et al. 2006]

$$\begin{pmatrix} r_{out} \\ g_{out} \\ b_{out} \end{pmatrix} = A_{3 \times 3} \begin{pmatrix} r_{in} \\ g_{in} \\ b_{in} \end{pmatrix} + B_{3 \times 1}$$



Output of Neural Style Transfer



Neural Style Transfer + Locally Affine Transformation

# Background on Neural Style Transfer

1. Decompose the input photo using a neural network

2. Impose some of the statistics of the model image



Preventing spillovers

3. Reconstruct the output image

# Original Algorithm and Its Limitation

- Transfer the Gram matrix of the model to the input

$$G_{ij}^{\ell} = \sum_k F_{ik}^{\ell} F_{jk}^{\ell}$$

$\ell$  layer index in the network  
 $i, j$  map index in the layer  
 $k$  pixel index in the map

# Original Algorithm and Its Limitation

- Transfer the Gram matrix of the model to the input

$$G_{ij}^\ell = \sum_k F_{ik}^\ell F_{jk}^\ell$$

*G does not depend on the pixel position*

$\ell$  layer index in the network

$i, j$  map index in the layer

$k$  pixel index in the map

# Original Algorithm and Its Limitation

- Transfer the Gram matrix of the model to the input

$$G_{ij}^{\ell} = \sum_k F_{ik}^{\ell} F_{jk}^{\ell}$$

*G does not depend on the pixel position*

$\ell$  layer index in the network

$i, j$  map index in the layer

$k$  pixel index in the map

- **Property:** the Gram matrix defines the vectors up to an isometry
  - “the neural responses can move but their distribution remains the same”
  - E.g., the G matrix implicitly encodes the “quantity of sky” in a photo.

# Different Horizon Lines



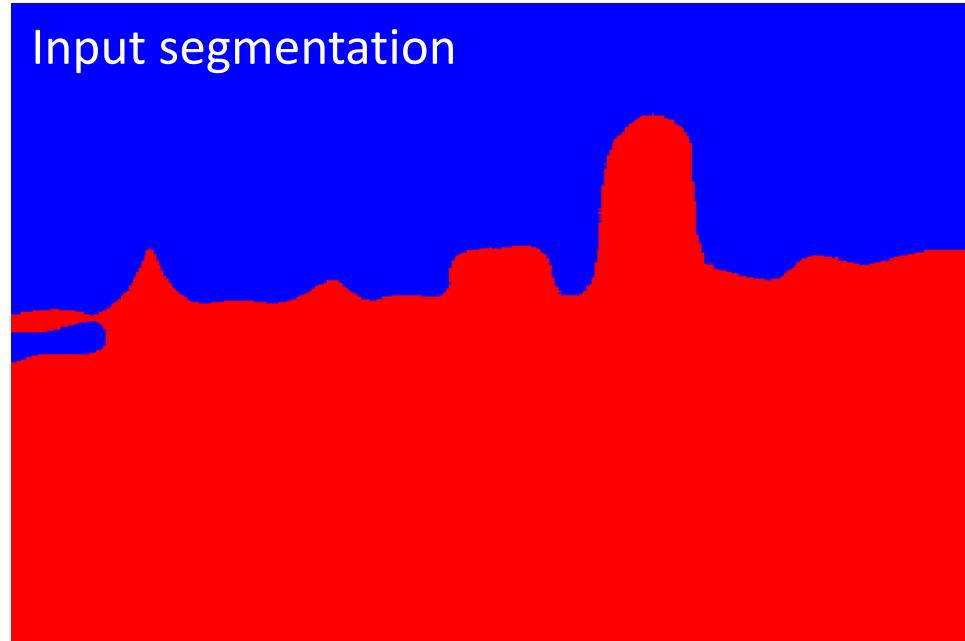
# Our Strategy to Ensure Consistent Matching

1. Run semantic segmentation on input and model [Chen et al. 2016]
  - Merge “visually equivalent” categories, e.g., sea and lake
  - Constrain the input labels to be a subset of the model labels
2. Transfer statistics within each category, e.g., sky to sky

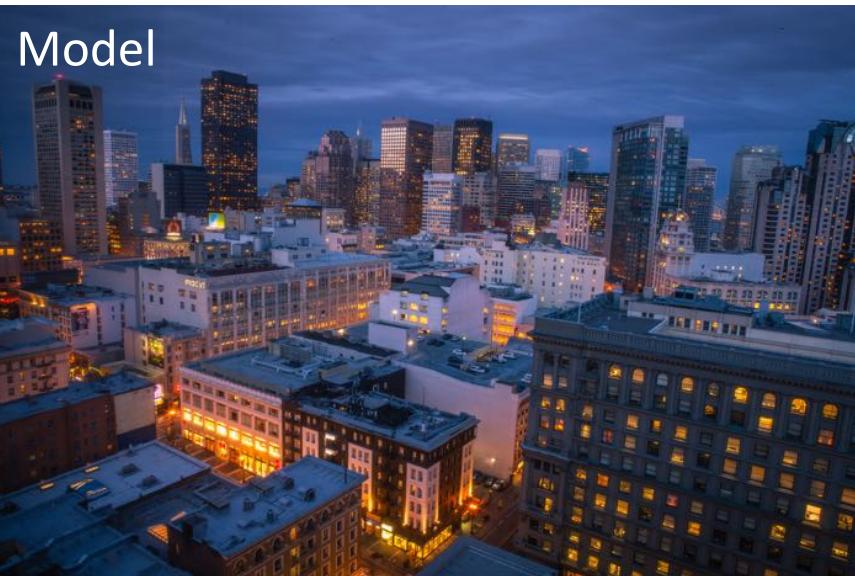
Input



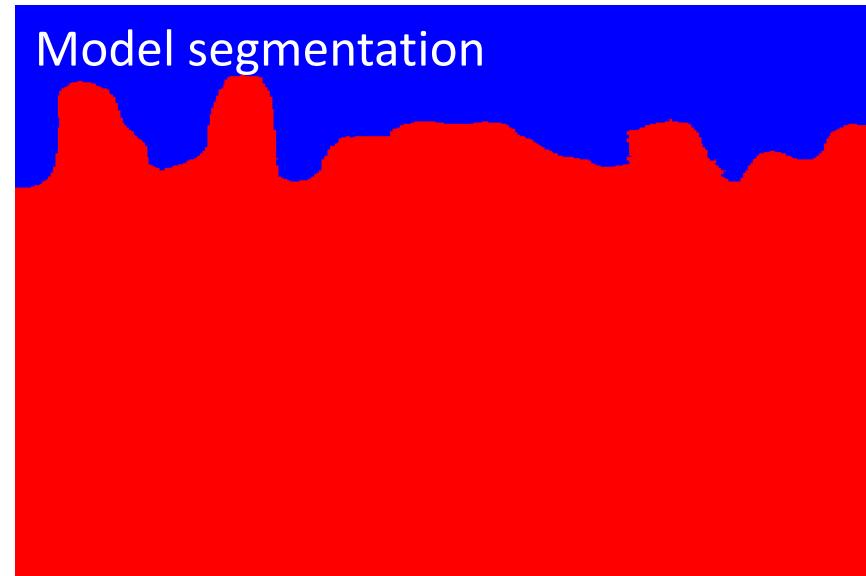
Input segmentation



Model



Model segmentation





Output of Neural Style Transfer



Model



Model

Output of Neural Style Transfer + Locally Affine Transformation



Model

Our result: Neural Style Transfer + Locally Affine Transformation + Semantic Matching



Input



Model



Model

Our result: Neural Style Transfer + Locally Affine Transformation + Semantic Matching



Input



Model





Neural Style Transfer



Model



Our result



Model



Input



Model



Neural Style Transfer



Model



Model

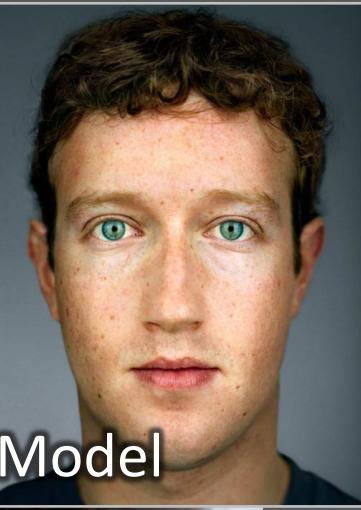
Color Histogram Transfer

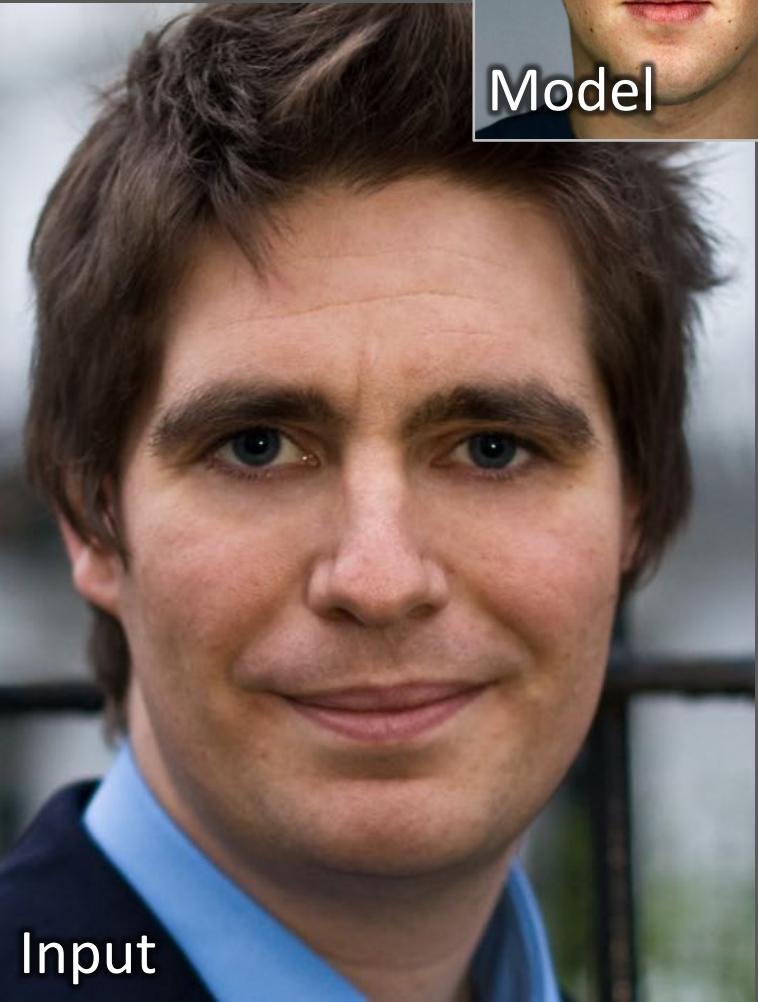


Our result

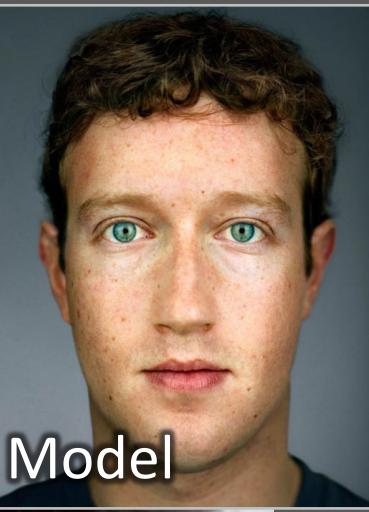


Model





Model



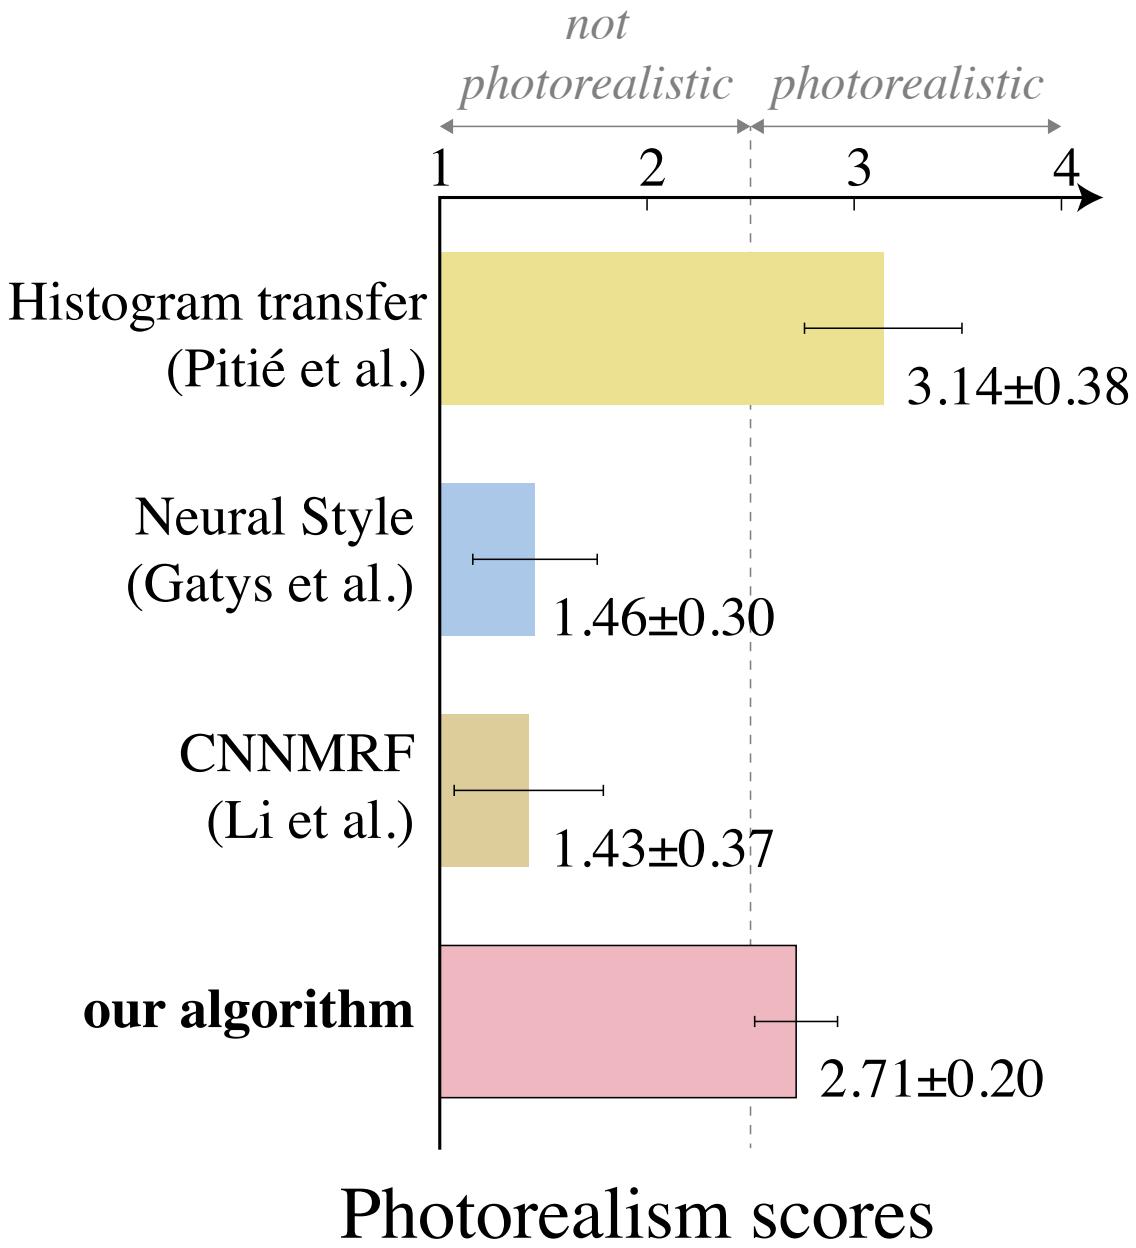
Input

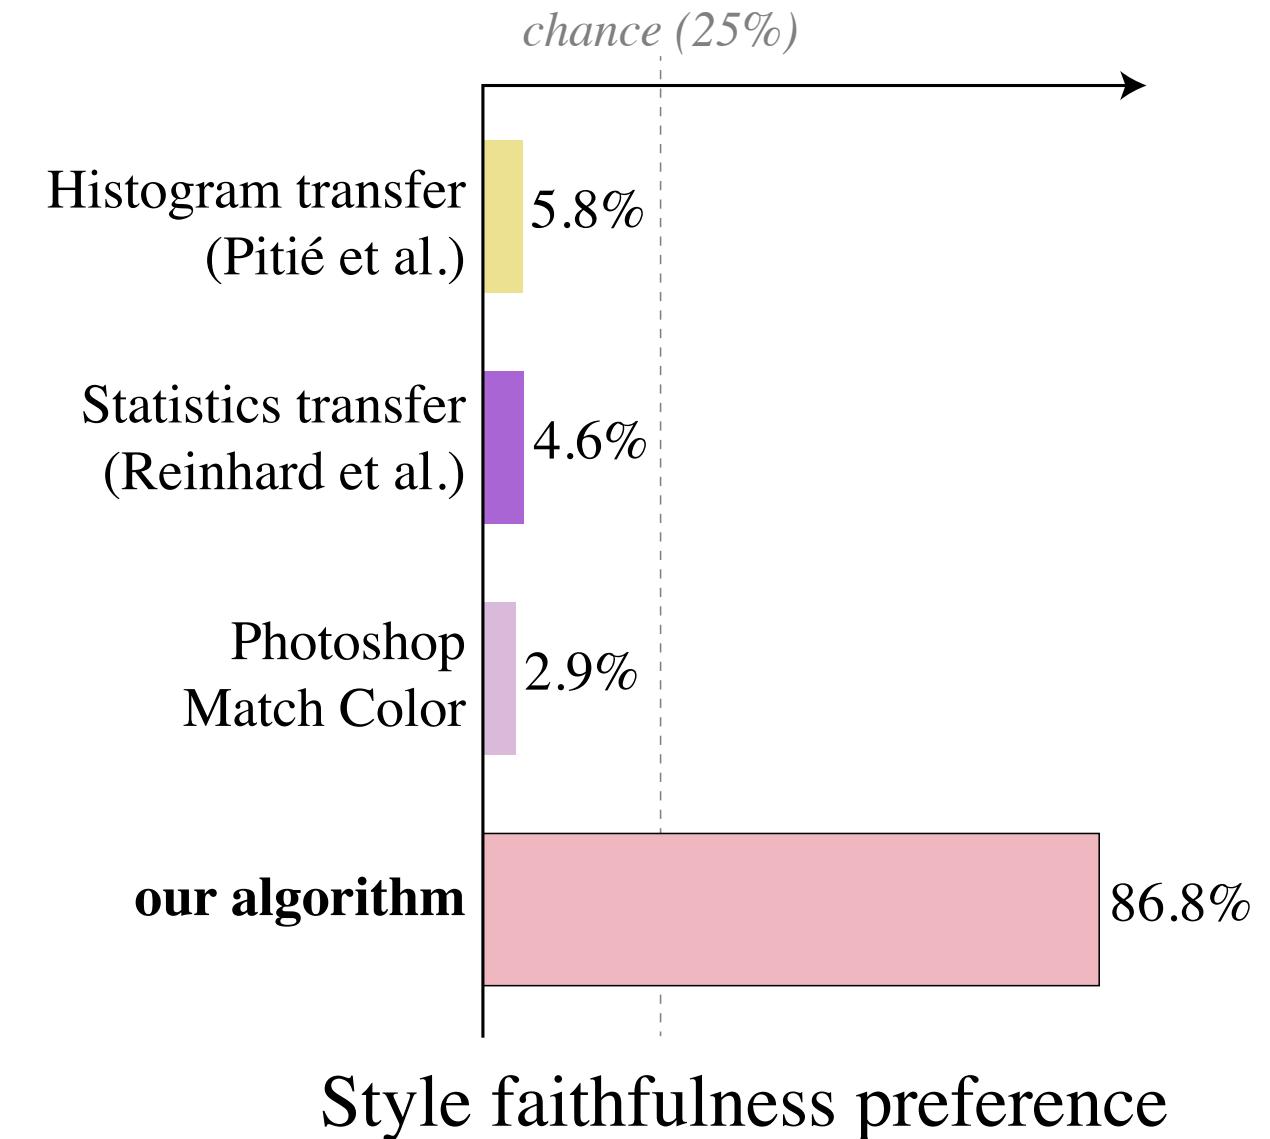
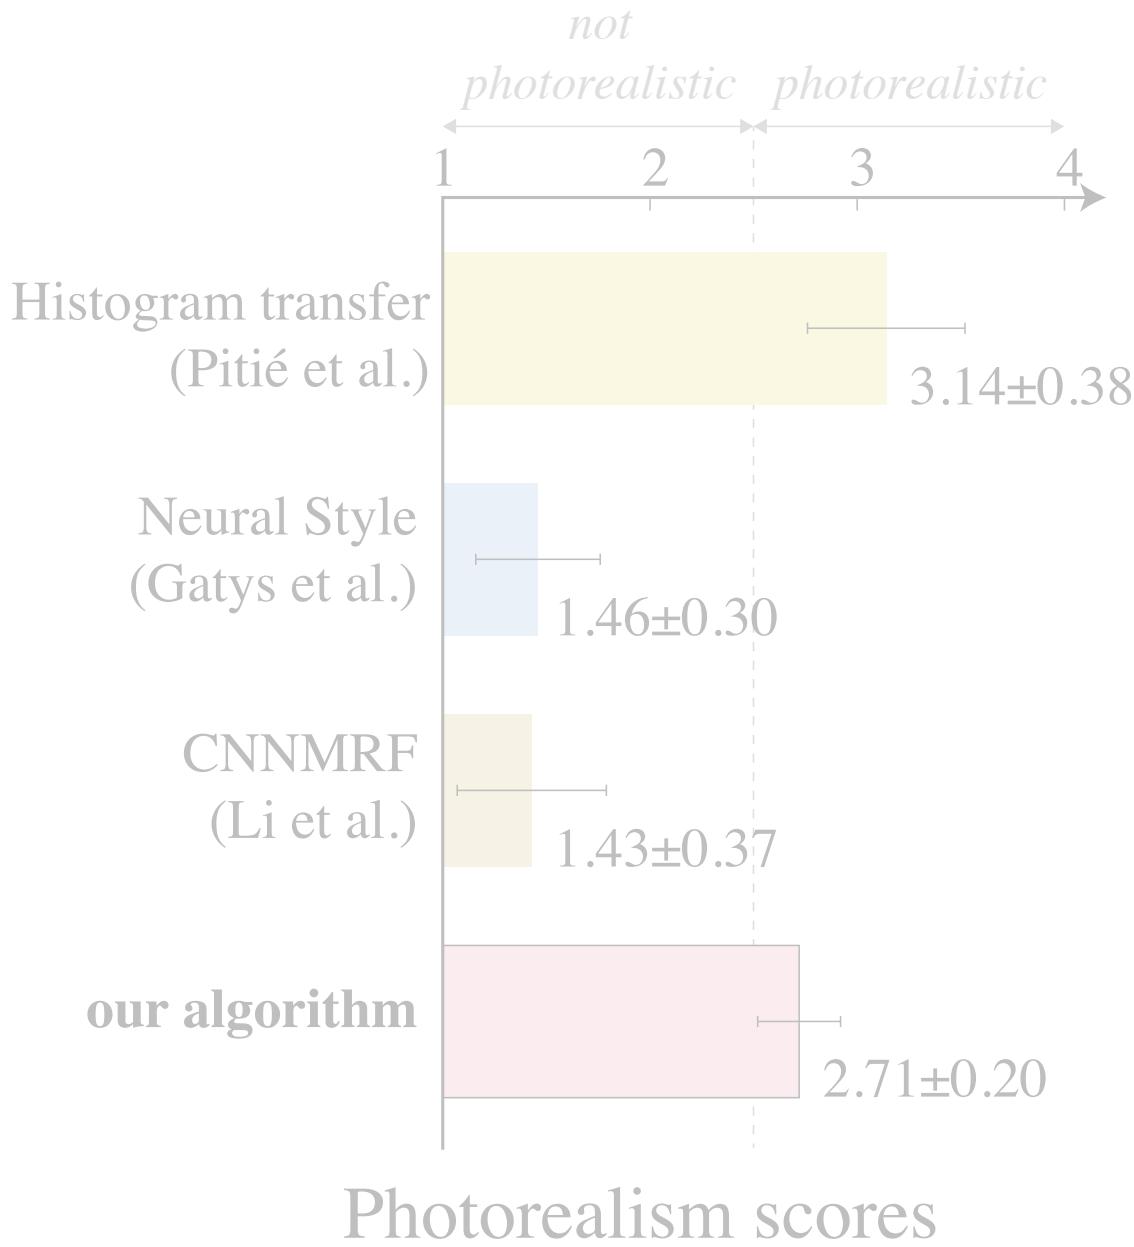


Our result



Portrait Style Transfer





*Disclaimer: specialized algorithms like to perform as well or better.*

# Required Steps for Productization

- Faster performance
  - Research prototype is too slow, on the order of minutes per photo
  - Feed-forward network, GPU, cloud...
- Smaller memory footprint
  - Currently based on VGG-19
  - Shallow network [Ustyuzhaninov et al., 2016]

What's next?

# Deep Photo Style Transfer is only a First Step

- Vast algorithm design space to explore
- Many questions to answer
  - Why a network trained for classification? Can we do better?
  - “Spillovers are obvious”, can we fix them without a full semantic analysis?
  - Where does the distortion come from?

## Background on Neural Style Transfer

1. Decompose the input and model images using a neural network
2. Impose some of the statistics of the model image
3. Reconstruct the output image

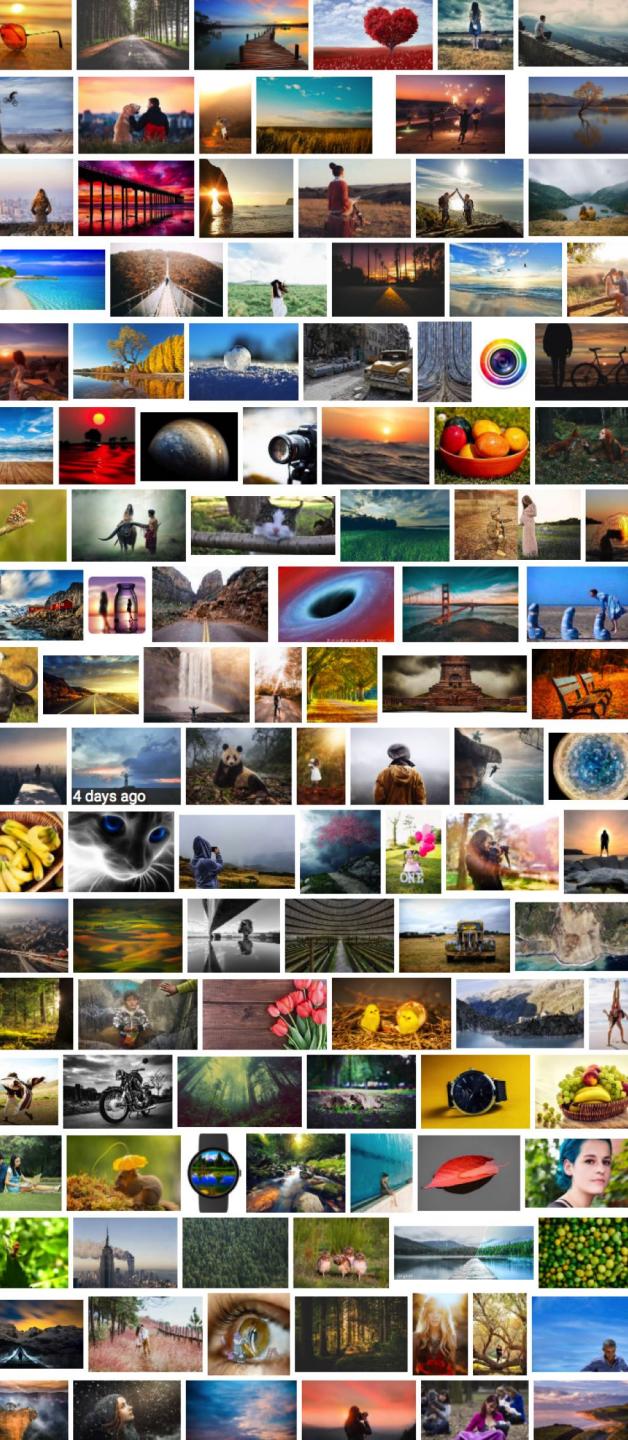
# Learning Creative Edits from Experts

- Challenge #1: experts cannot create 1 million+ photos for training a CNN
- Challenge #2: the range of effects is large
- Challenge #3: not all styles apply to all photos



# Exploiting On-line Collections

- What can we do with 1 million or more photos?
- How do we find the “good ones” and exploit them?
- Working with sets instead of samples à la CycleGAN or Exposure is promising but still needs work  
[Zhu et al. 2017, Zhang et al. 2018]



# A Great Photo from Every Shot



# Better Photos, Easily



# Photo Style Transfer

- A simple and powerful way to edit photos
- Challenges us to better understand images
  - What is style? What is content?
  - What makes an image look like it does?
  - How to preserve the image structure?
- What I did not show: videos, panoramas, weather, lighting, perception, performance...



# Thanks!

Nothing would have been possible without my collaborators:  
Mathieu Aubry, Soonmin Bae, Kavita Bala, Connally Barnes,  
Adrien Bousseau, Vladimir Bychkovsky, Eric Chan, Jiawen Chen,  
Jeff Chien, George Drettakis, Frédo Durand, Bill Freeman, Sam Hasinoff,  
Jan Kautz, Fujun Luan, Eli Shechtman, and Yichang Shih